

Learning State Machines

Learning in the limit, Myhill-Nerode, the Hankel matrix,
Active and Passive learning

Learning (reverse-engineering)

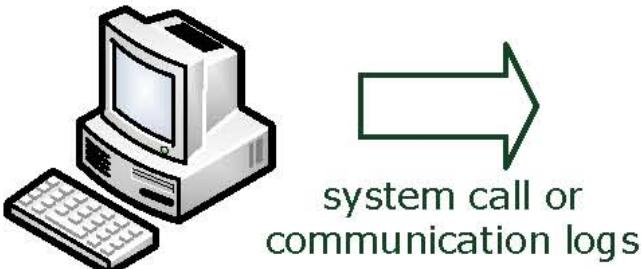
- Execution logs form a **big data** set from which can be used to gain information about a system or protocol

This can help to

- analyze **your own code** and hunt for bugs, or
- reverse-engineer **someone else's unknown protocol**, eg. a botnet, to fingerprint or to analyze (and attack)

State machine learning

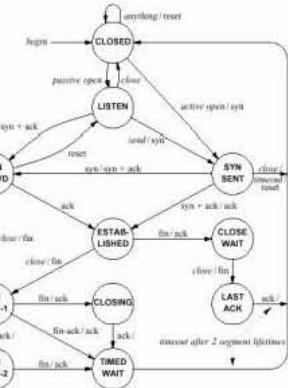
software system



execution traces

A B E E E D ...
A B D C D C D ...
B B B D G H A ...
B B D D D E H ...
...

state machine learning



software model

- Software leaves **traces**
- A state machine is a logical **model** describing these traces
 - Classification – is a new trace generated by the same software?
 - Prediction – what trace is most likely to occur next?
 - **Analysis** – is the software deadlock-free, secure, **malicious**?

Learning DFAs from labeled data

4

- Input:
 - Labeled data, positive and negative traces
- Goal:
 - Find the **smallest** DFA that is **consistent** with the data

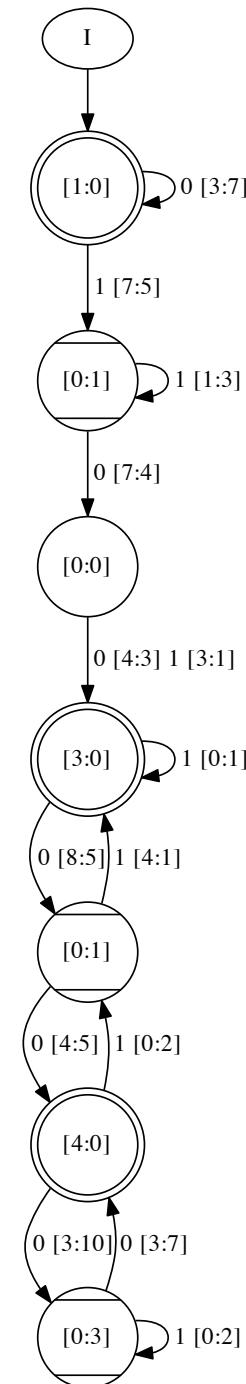
Learning DFAs from labeled data

5

- Input:
 - Labeled data, positive and negative traces
- Goal:
 - Find the **smallest** DFA that is **consistent** with the data
- Occam's razor
 - *Among competing hypotheses, the one with the fewest assumptions should be selected.*

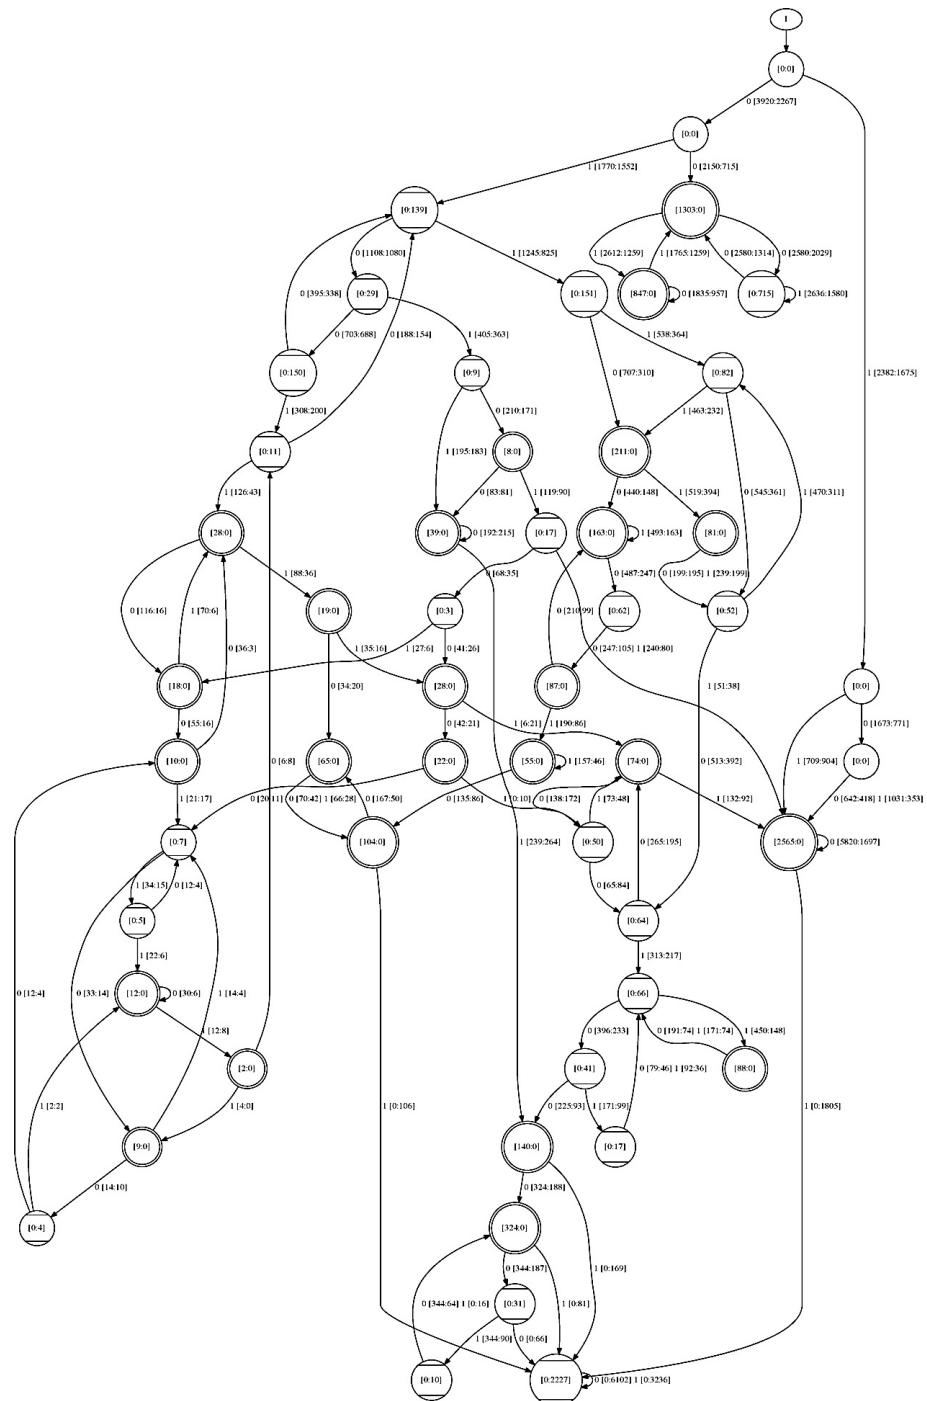
Example

```
19100000000  
115011000101010100  
0501000  
1200  
012110000000100  
1510100  
13100  
06000001  
13101  
0201000001010000011100  
0  
0130111010000000  
13101  
171000000
```

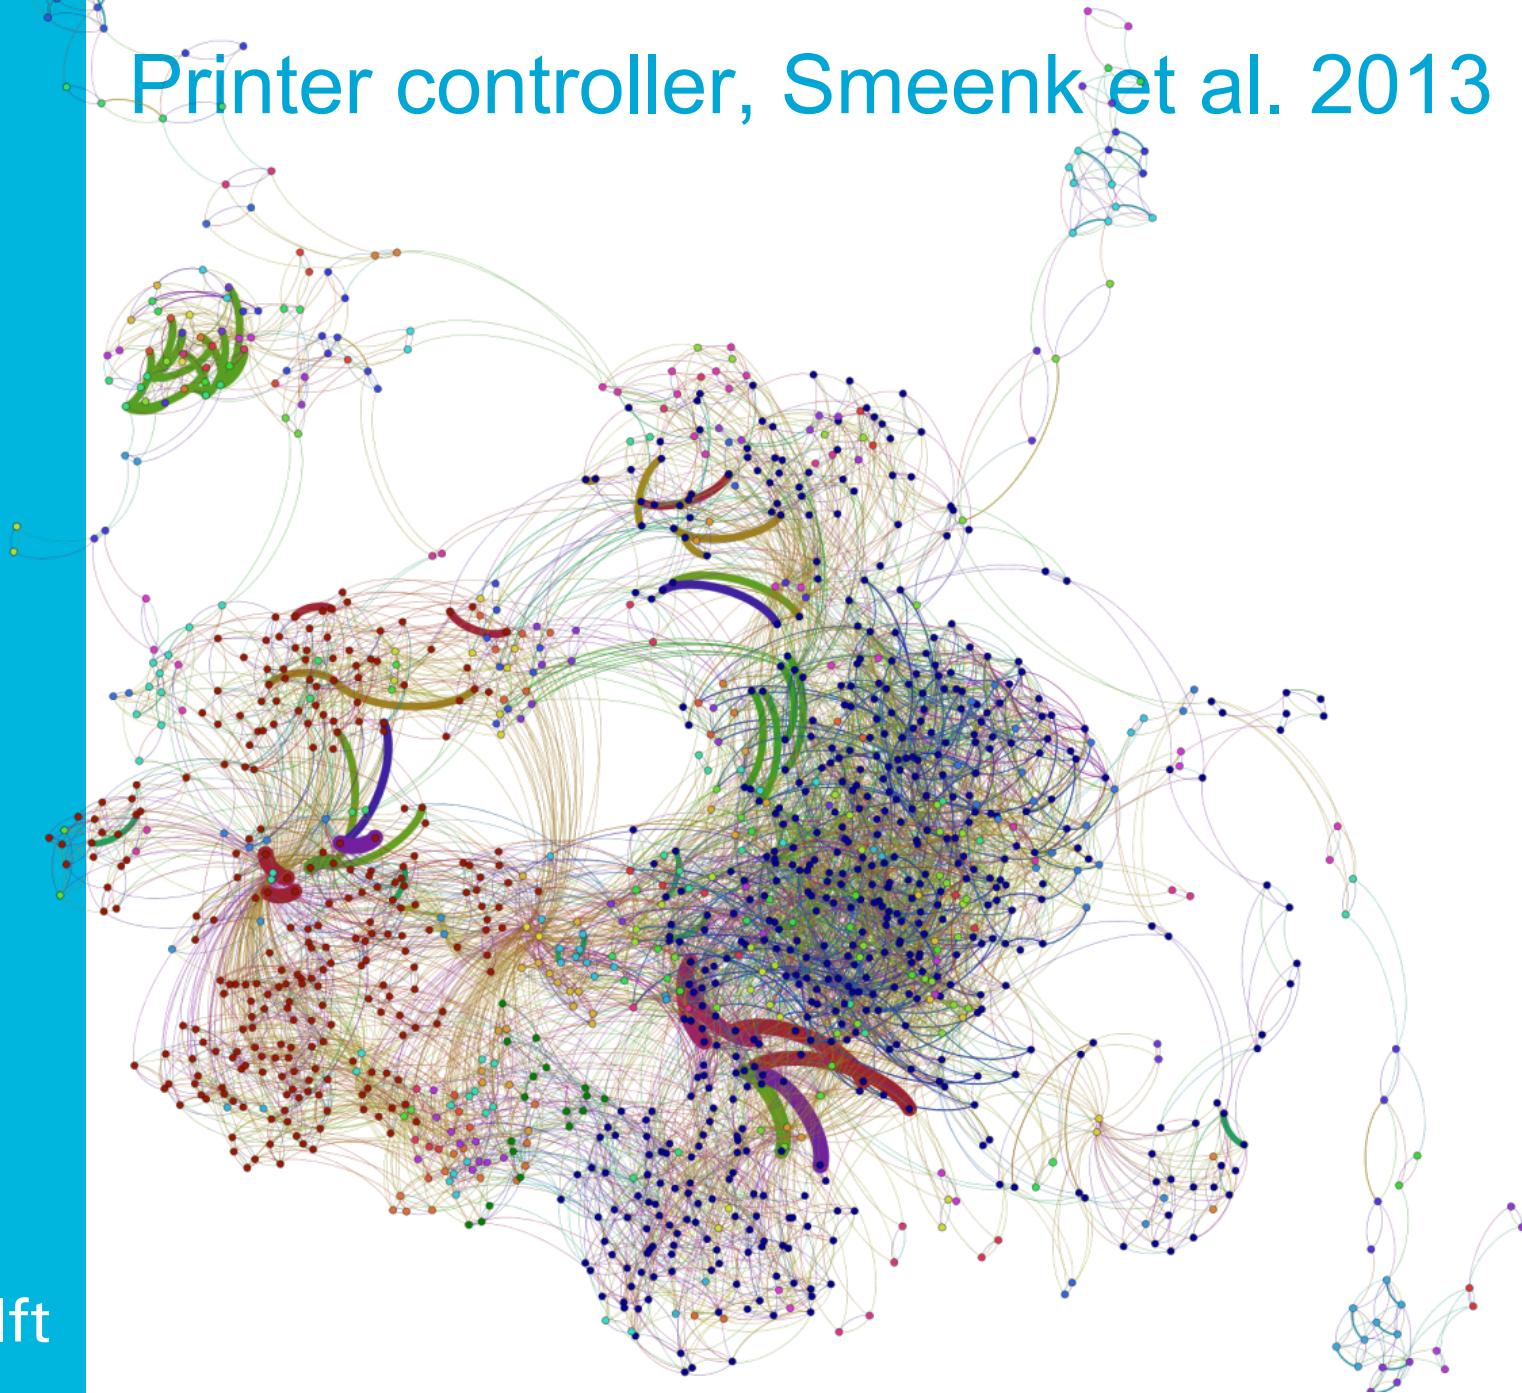


Example

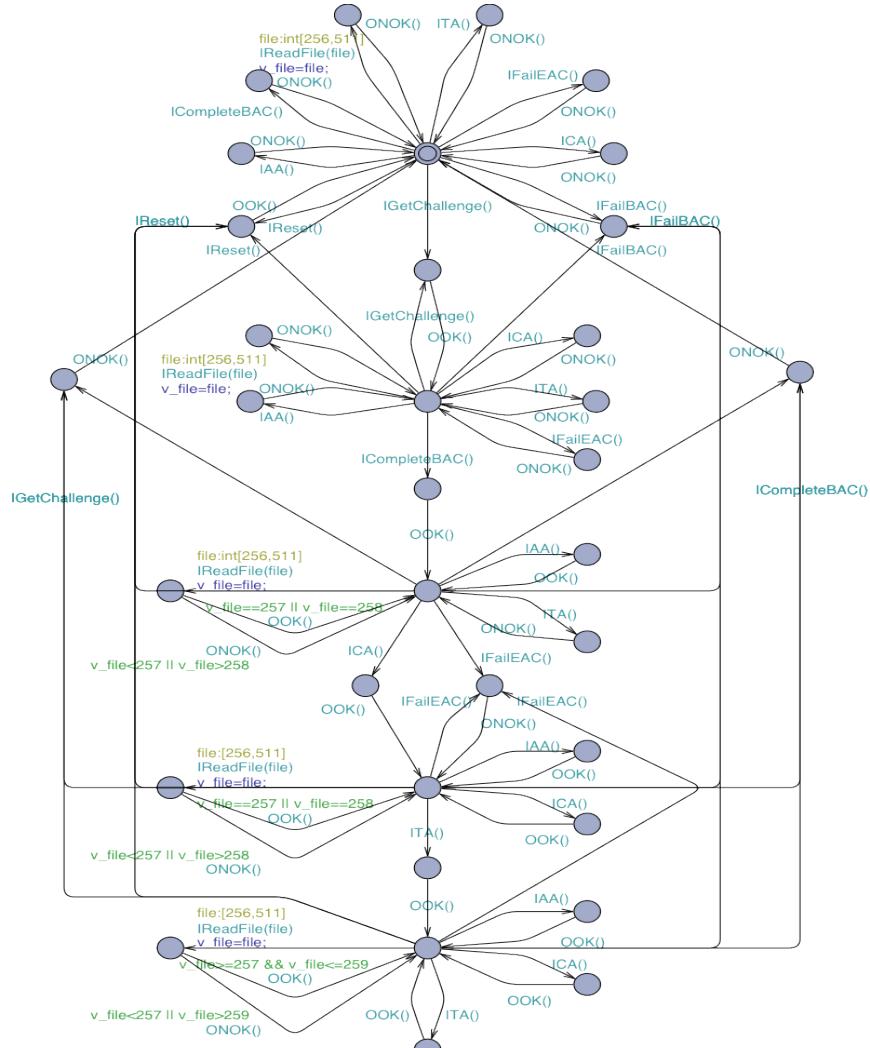
19100000000
115011000101010100
0501000
1200
012110000000100
1510100
13100
06000001
13101
02010000010100000111000
0130111010000000
13101
171000000
...
10244 strings in total



Printer controller, Smeenk et al. 2013



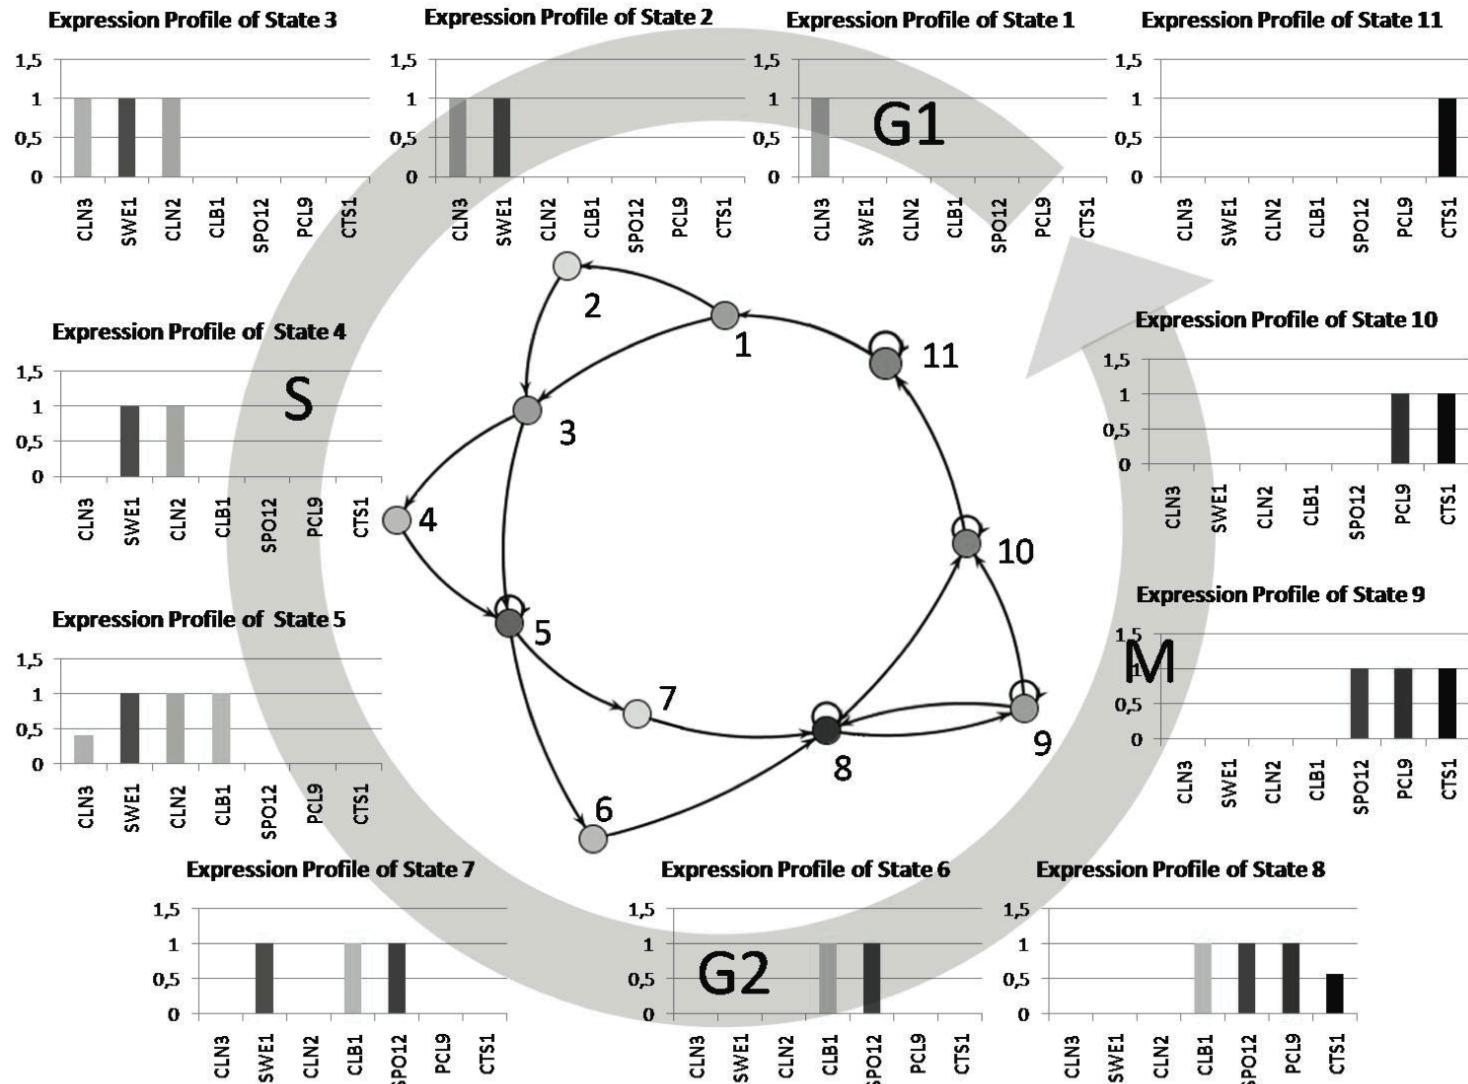
Biometric Passport, Aarts et al. 2010



Using state machines

- Analyzing the models by hand, or with model checker, for flaws
 - to see if *all paths* are correct & secure
- Fuzzing or model-based testing
 - using the diagram as basis for “deeper” fuzz testing
- Program verification
 - *proving* that there is no functionality beyond that in the diagram, which using testing you can never establish
- Using it when doing a manual code review

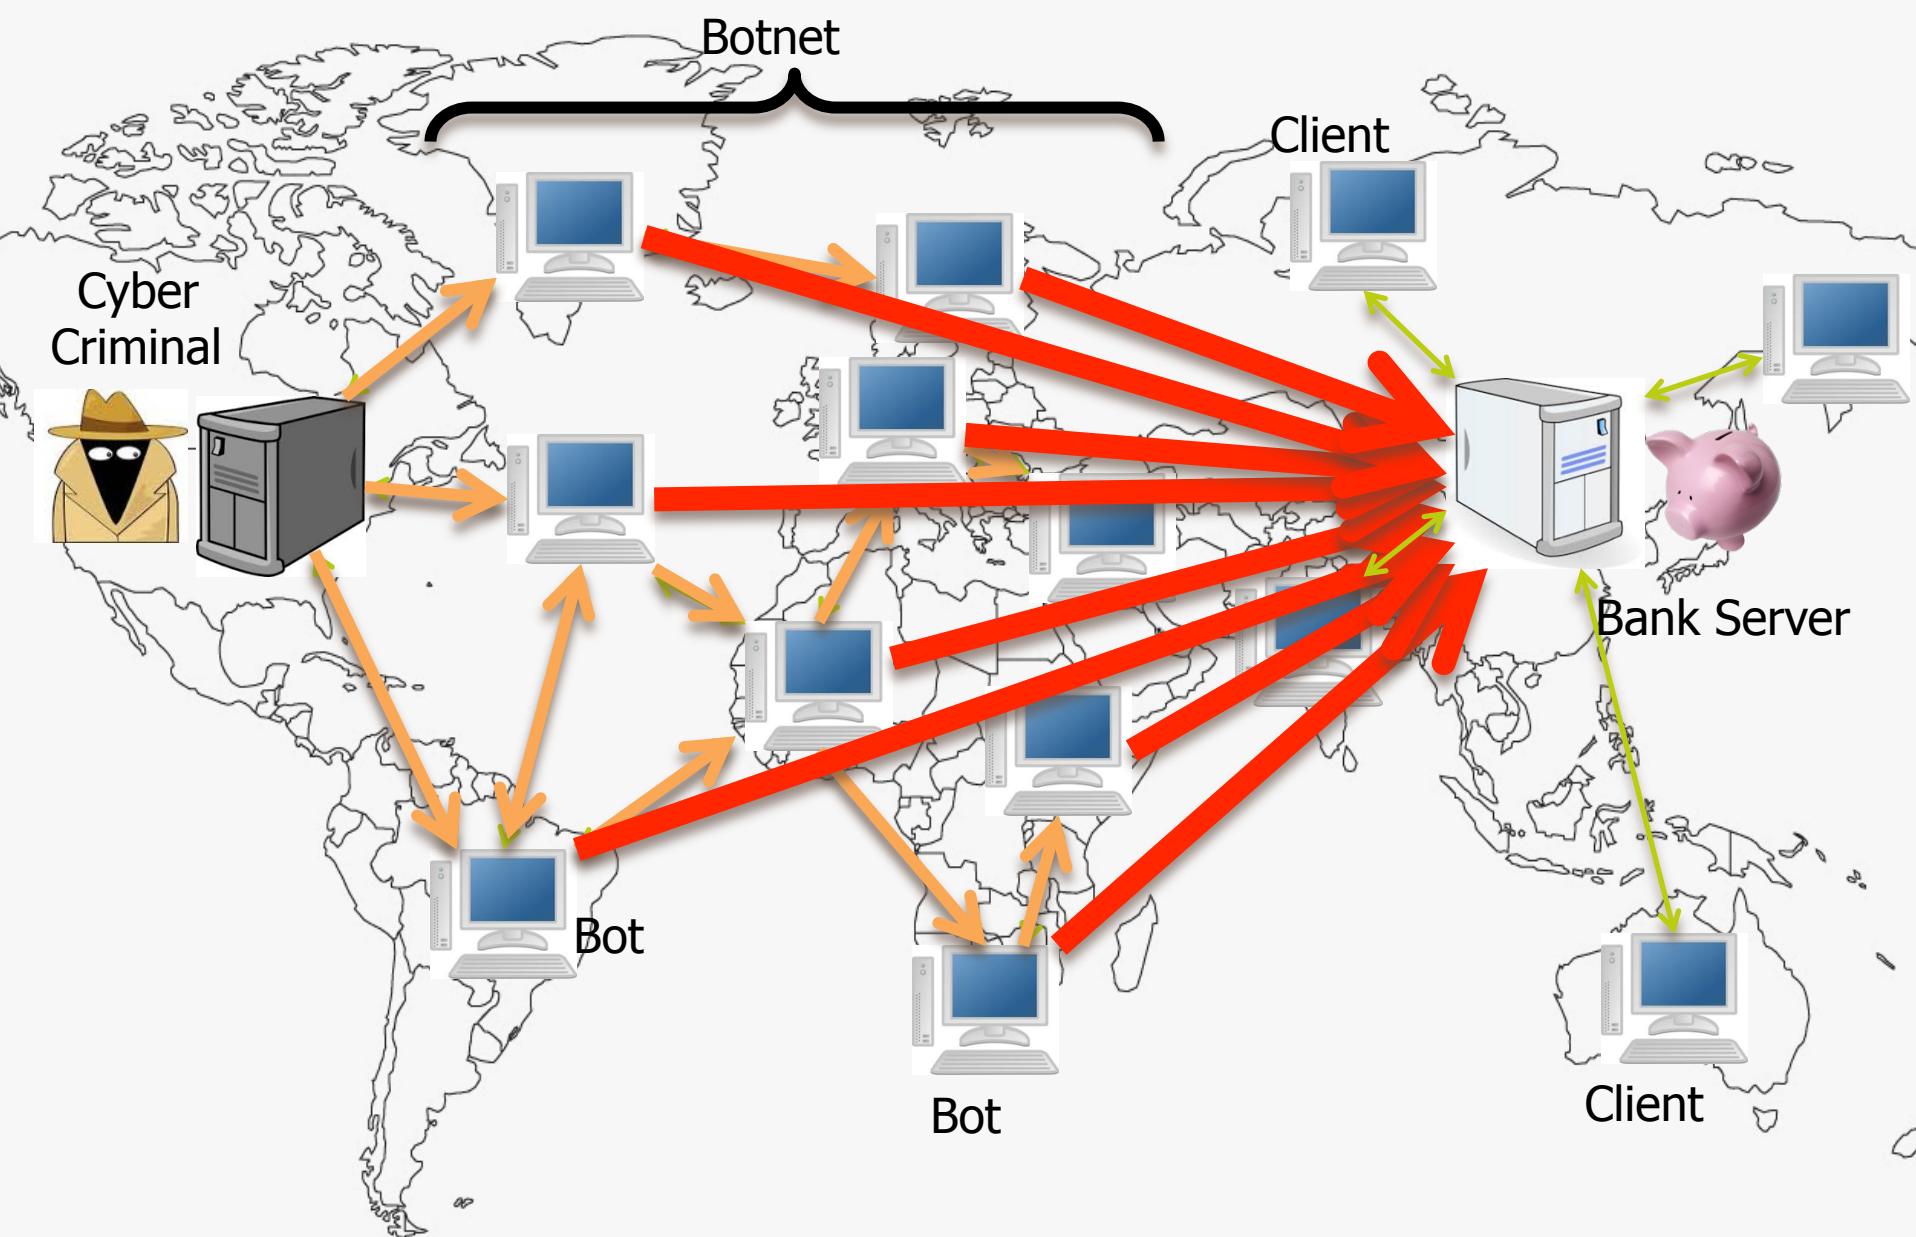
Yeast cycle, Schmidt et al. 2013



Preventing botnets

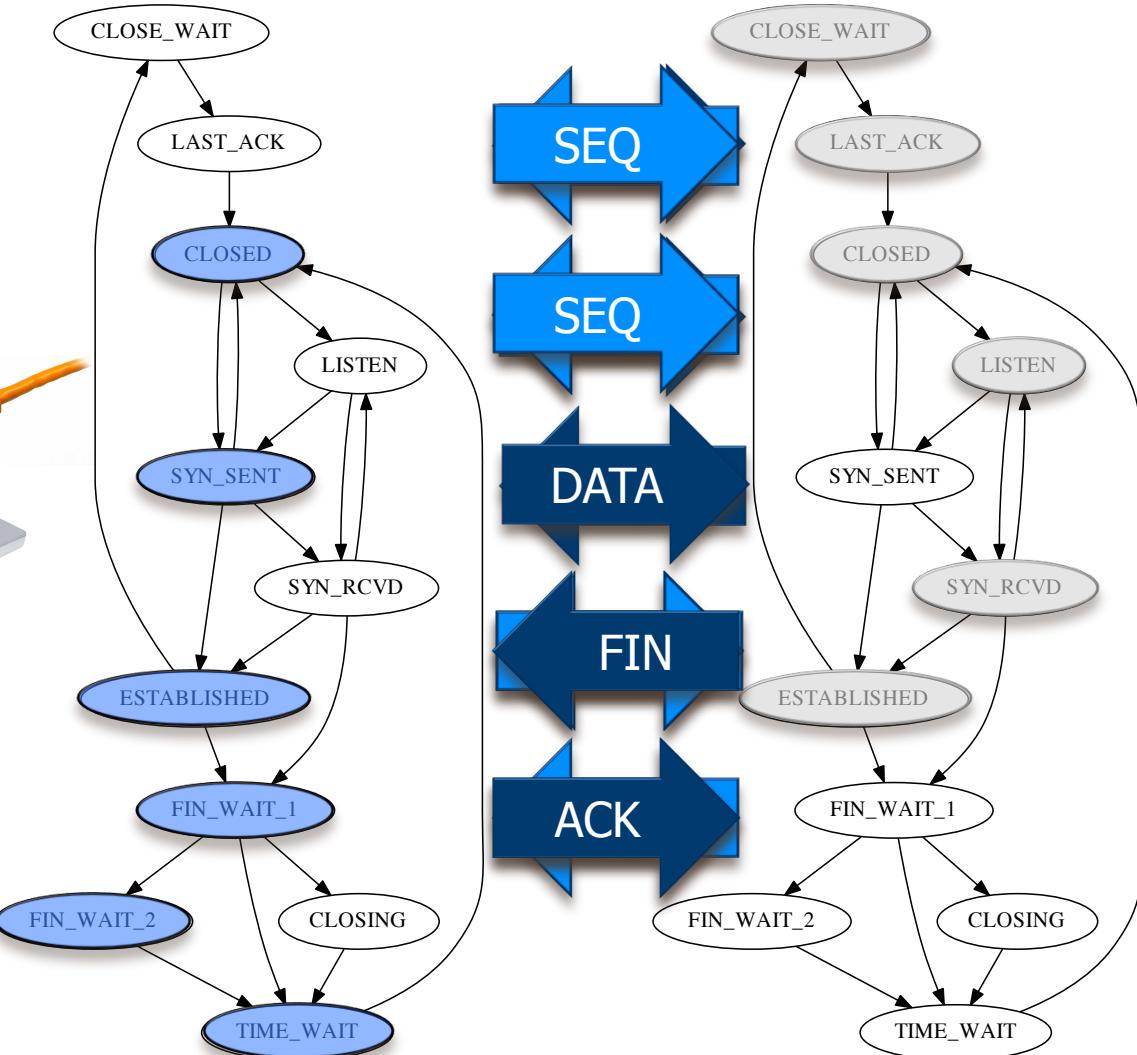


- Traditional
 - code analysis and finding malware fingerprints
- Code/binary analysis is mostly **manual** and increasingly **harder**
 - Code obfuscation
 - Encryption
 - Self-modifying
- Behavior-based analysis is much harder to thwart
 - Bots need to **communicate!**



We observe internet traffic

- SYN, SYN-ACK, ACK, SEQ, ACK, SEQ, ACK, SEQ, ACK, FIN-ACK, ACK, FIN-ACK
- SYN, SYN-ACK, ACK, FIN-ACK, ACK, FIN-ACK
- SYN, SYN-ACK, ACK, SEQ, ACK, FIN-ACK, ACK, FIN-ACK
-



Learning (reverse-engineering)

- Execution logs form a **big data** set from which can be used to gain information about a system or protocol

This can help to

- analyze **your own code** and hunt for bugs, or
- reverse-engineer **someone else's unknown protocol**, eg. a botnet, to fingerprint or to analyze (and attack)

We observe internet traffic

- SYN, SYN-ACK, ACK, SEQ, ACK, SEQ, ACK, SEQ, ACK, FIN-ACK, ACK, FIN-ACK
- SYN, SYN-ACK, ACK, FIN-ACK, ACK, FIN-ACK
- SYN, SYN-ACK, ACK, SEQ, ACK, FIN-ACK, ACK, FIN-ACK
-

First: make abstract

- A, B, C, D, C, D, C, D, C, E, C, E
- A, B, C, E, C, E
- A, B, C, D, C, E, C, E
-

Then: look for patterns!

- A, B, C, D, C, D, C, D, C, E, C, E
- A, B, C, E, C, E
- A, B, C, D, C, E, C, E
-

Then: look for patterns!

- A, B, C, D, C, D, C, D, C, E, C, E
- A, B, C, E, C, E
- A, B, C, D, C, E, C, E
-

Then: look for patterns!

- A, B, C, D, C, D, C, D, C, E, C, E
- A, B, C, E, C, E
- A, B, C, D, C, E, C, E
-

Then: look for patterns!

- A, B, C, D, C, D, C, D, C, E, C, E
- A, B, C, E, C, E
- A, B, C, D, C, E, C, E
-

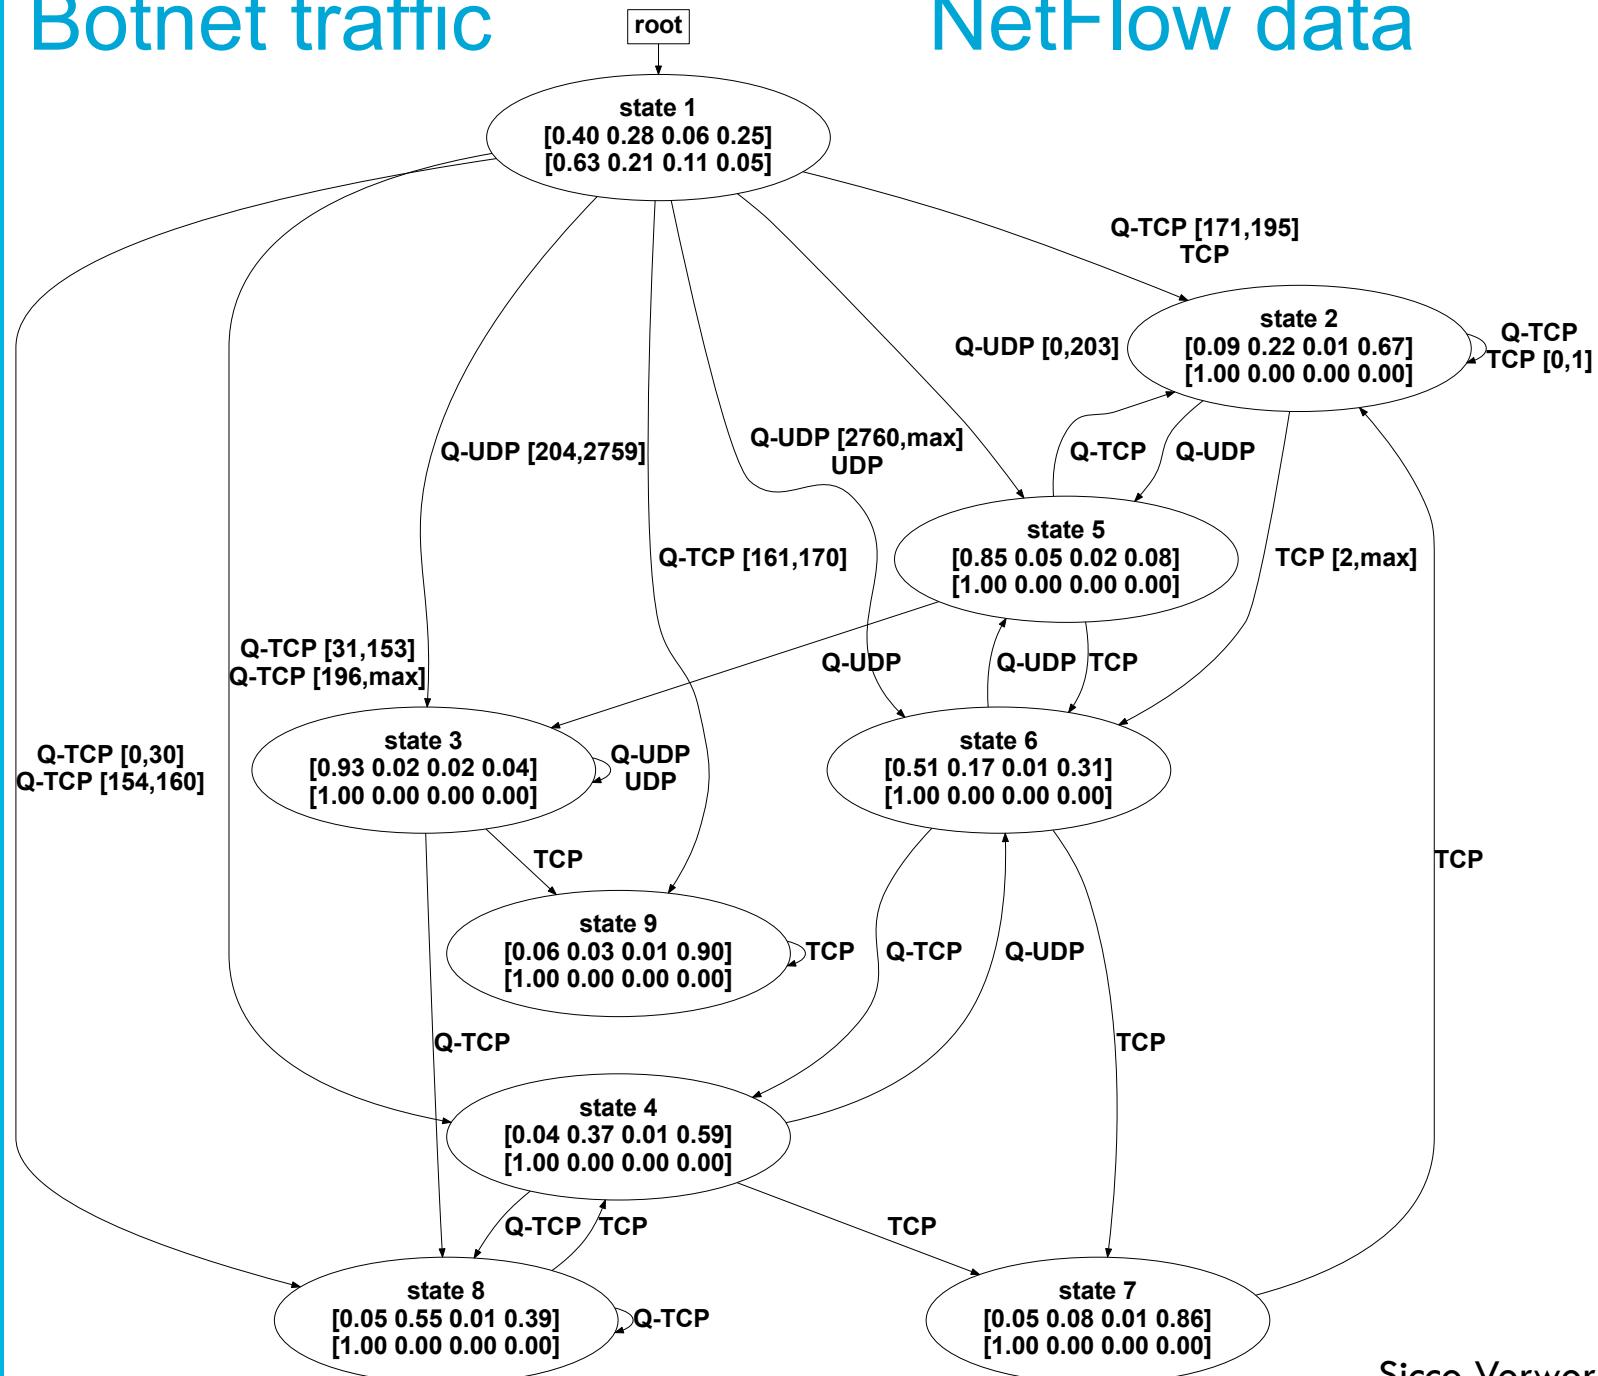
Then: look for patterns!

- A, B, C, D, C, D, C, D, C, E, C, E
- A, B, C, E, C, E
- A, B, C, D, C, E, C, E
-



Botnet traffic

NetFlow data



NetFlow data --- fingerprinting

- 4 scenarios: click-fraud, spamming, DDoS, ...
- For each scenario:
 - Learn state machine model from 1 malicious host
 - Behavioral difference is small → malicious

TP	TN	FP	FN
9	377	0	0
9	153	0	0
1	35	0	1
0	7	0	2

NetFlow data --- anomaly detection

- 4 scenarios: click-fraud, spamming, DDoS, ...
- For each scenario:
 - Learn state machine model from a large set of benign hosts
 - Behavioral difference to every model is large → malicious

TP	TN	FP	FN
10	561	3	0
10	225	6	0
2	50	4	1
0	8	1	3

WHAT IS LEARNING?

Computational complexity

- NP
 - Class of problems **verifiable** in **polynomial time**
- NP-hard
 - **At least as difficult to solve as any other problem in NP**
 - Solving in polynomial time is widely believed to be impossible (even with quantum computers)
- NP-complete
 - NP-hard and in NP, e.g., satisfiability, traveling salesman
- Undecidable
 - Impossible to solve using a computer, equivalent to the halting problem

Learning theory

- Field that studies **learning problems**, or how to find a model for data
 - computational complexity
 - approximation algorithms
 - statistical guarantees
 - convergence properties
 - ...
- One of the oldest fields in computer science, contributed to many innovations
 - neural networks, support vector machines, minimum description length
 - grammatical inference, state machine learning

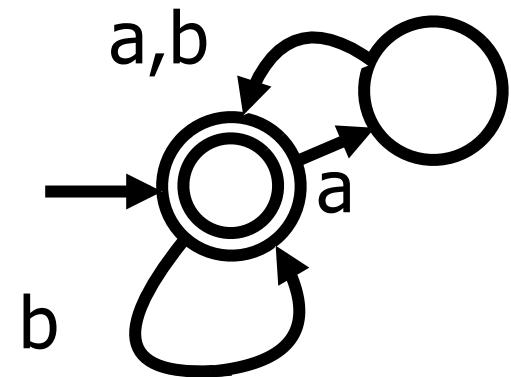
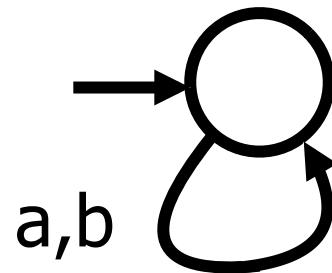
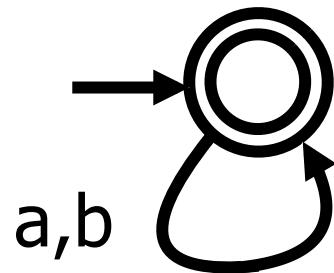
Learning in the limit

- Studies learning from **infinite data**
- Views learning as an ongoing process
 1. learner **receives some data** from a target concept C
 2. **updates** its hypothesis H, then goes to 1.
- Learning in the limit is successful if it **converges**:
 - $L(H) = L(C)$, at some point in time t
 - and $L(H) = L(C)$ at all time points $> t$
 - All data will be presented to the learner
 - Learner does not need to know when it converges

Learning by enumeration

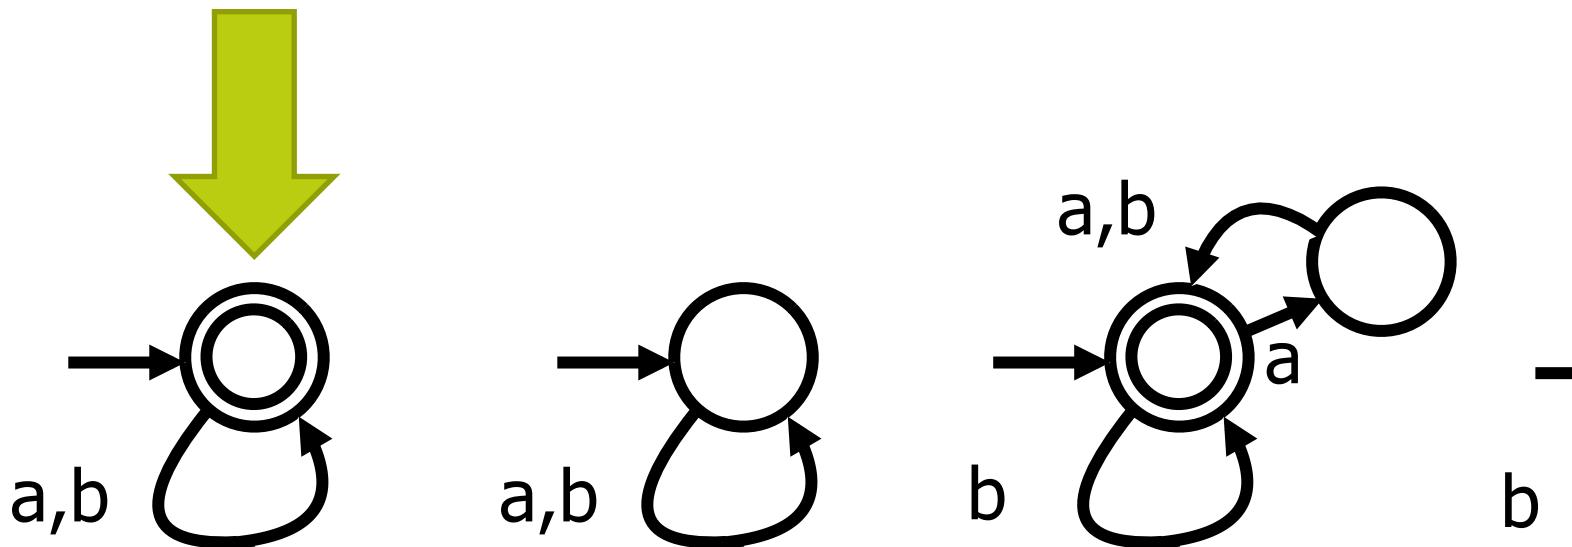
- A simple learning algorithm:
 - Given an **enumeration** H_1, H_2, H_3, \dots
 - Assume H_1 until it is inconsistent with the data
 - Then assume the first H_n such that there is no H_m with $m < n$ that is consistent with the data
 - Read more data and continue
- Many learning algorithms are of this type, but traverse the enumeration in different (smarter) ways

Learning by enumeration



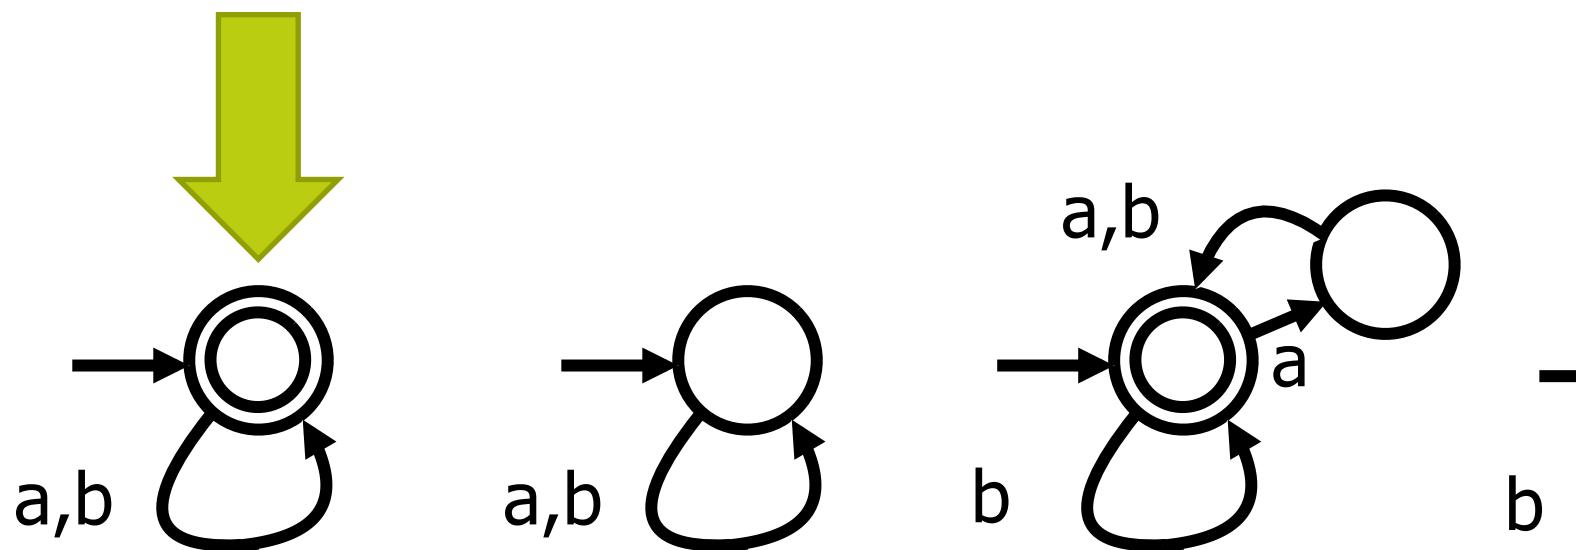
a
—
b

Learning by enumeration



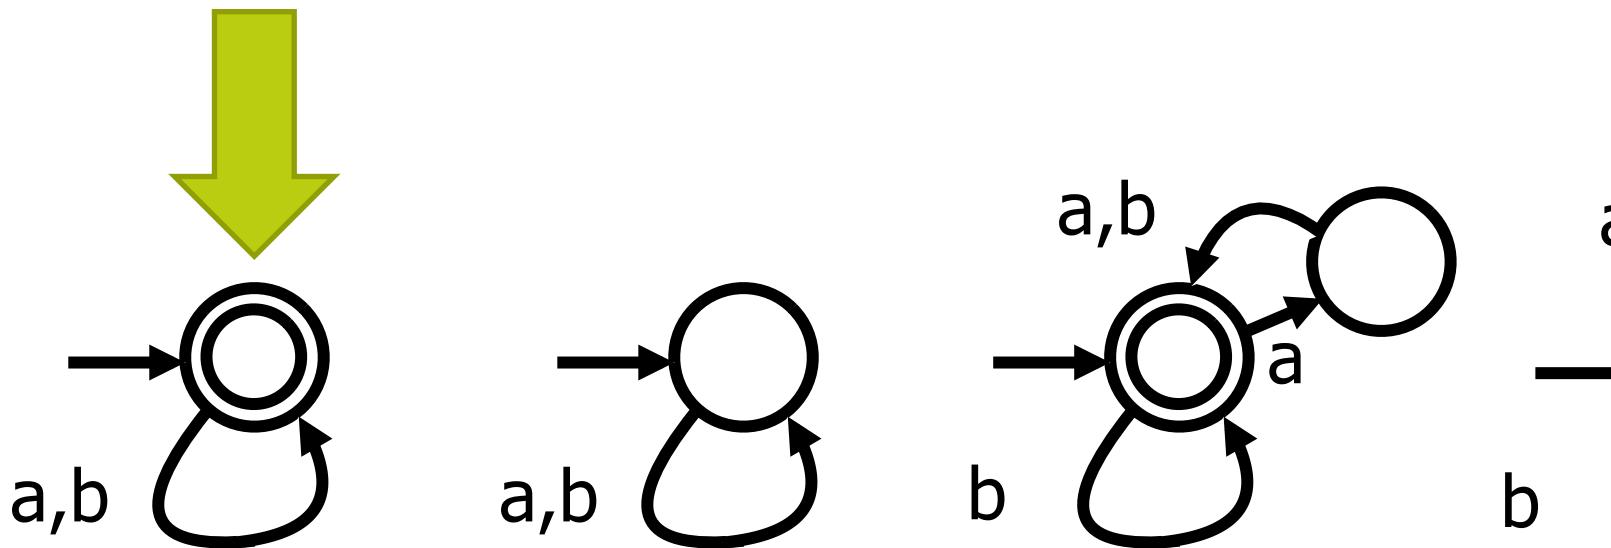
Examples: (a, +)

Learning by enumeration



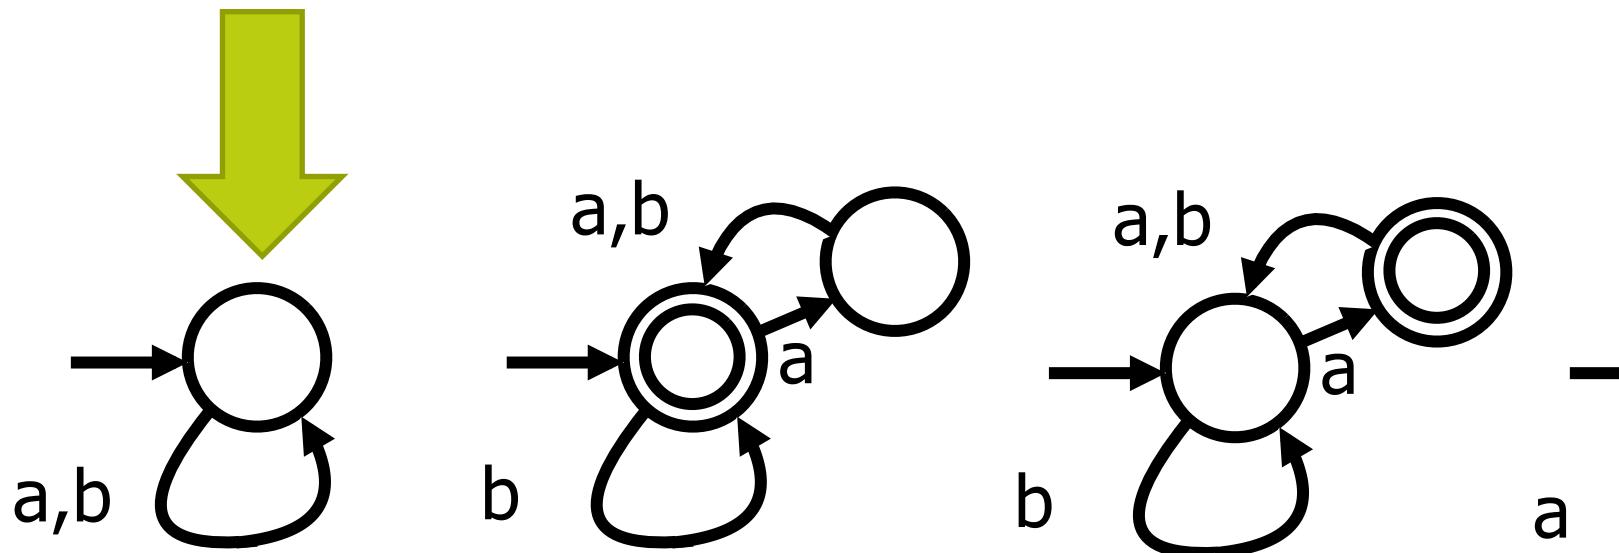
Examples: $(a,+)(aa,+)$

Learning by enumeration



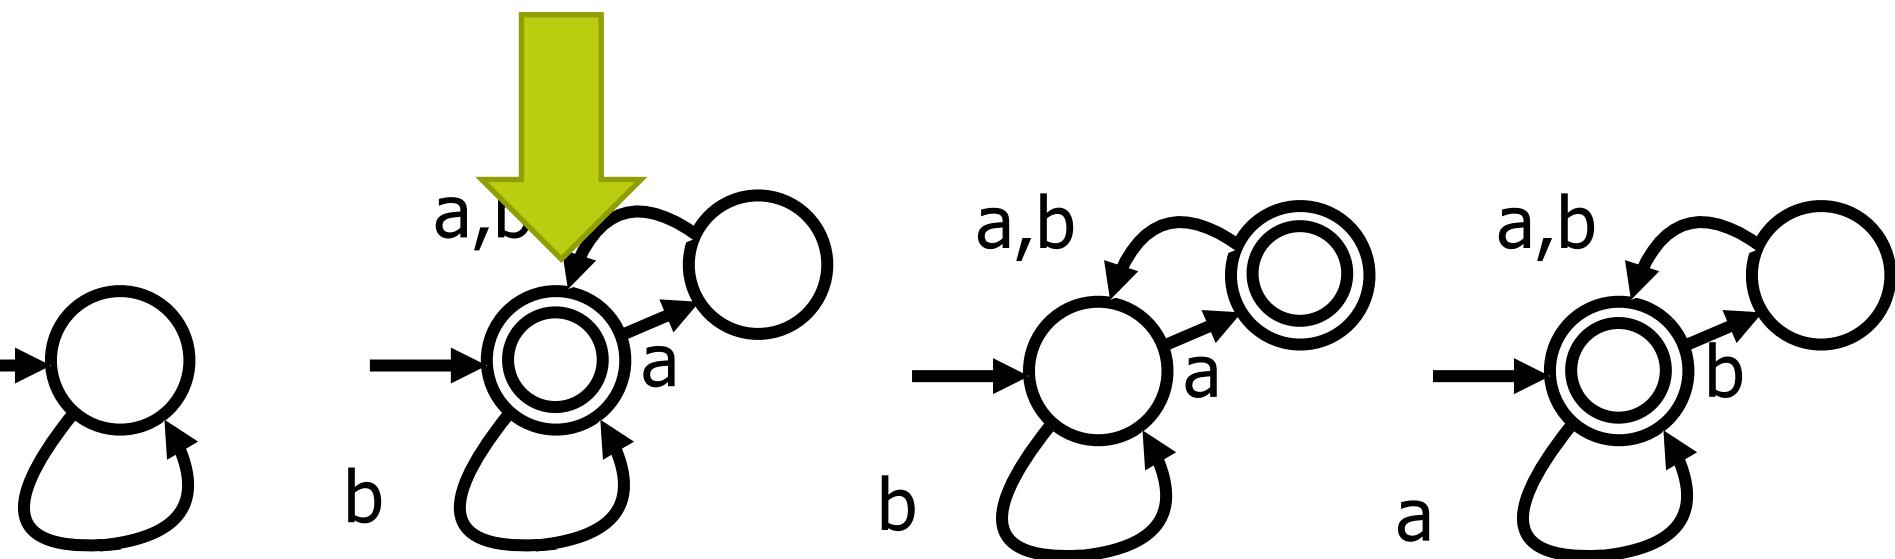
Examples: $(a,+)(aa,+)(b,-)$

Learning by enumeration



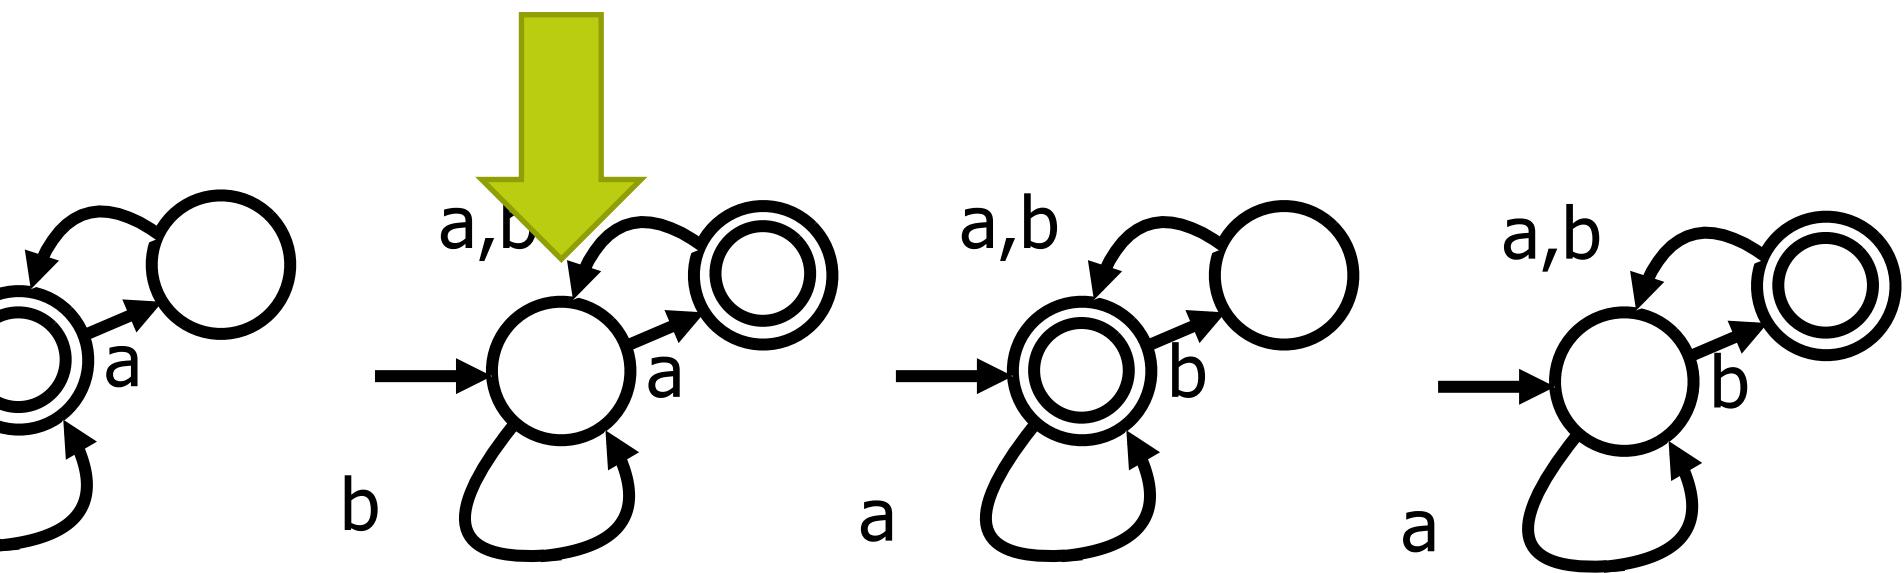
Examples: $(a,+)(aa,+)(b,-)$

Learning by enumeration



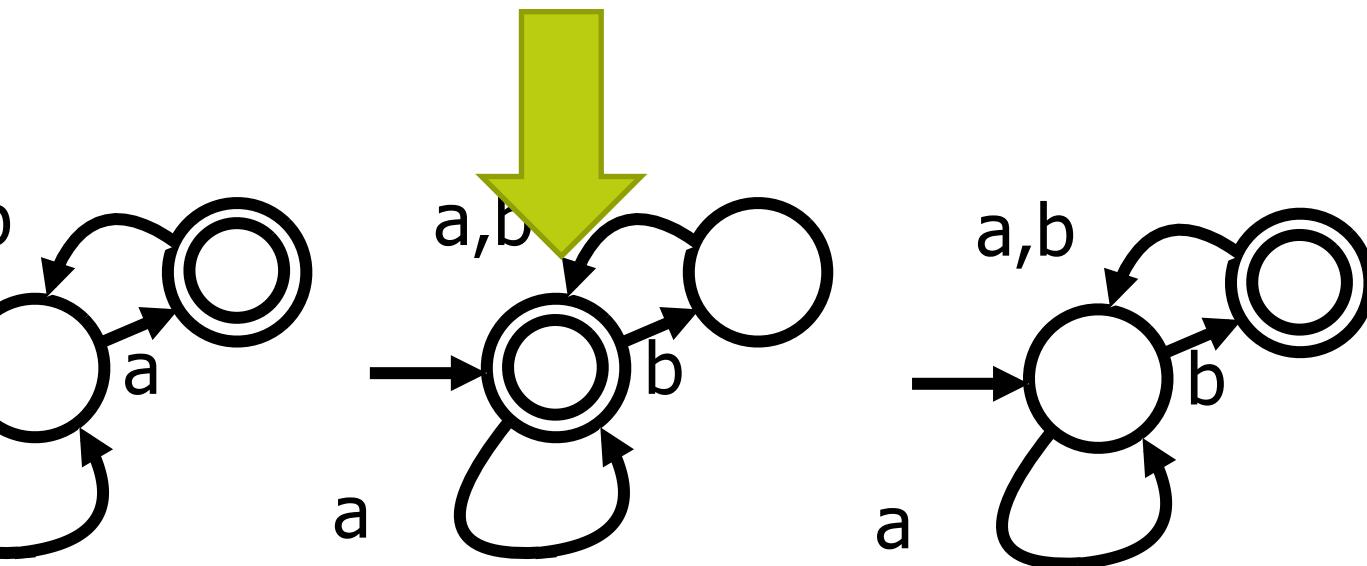
Examples: $(a,+)(aa,+)(b,-)$

Learning by enumeration



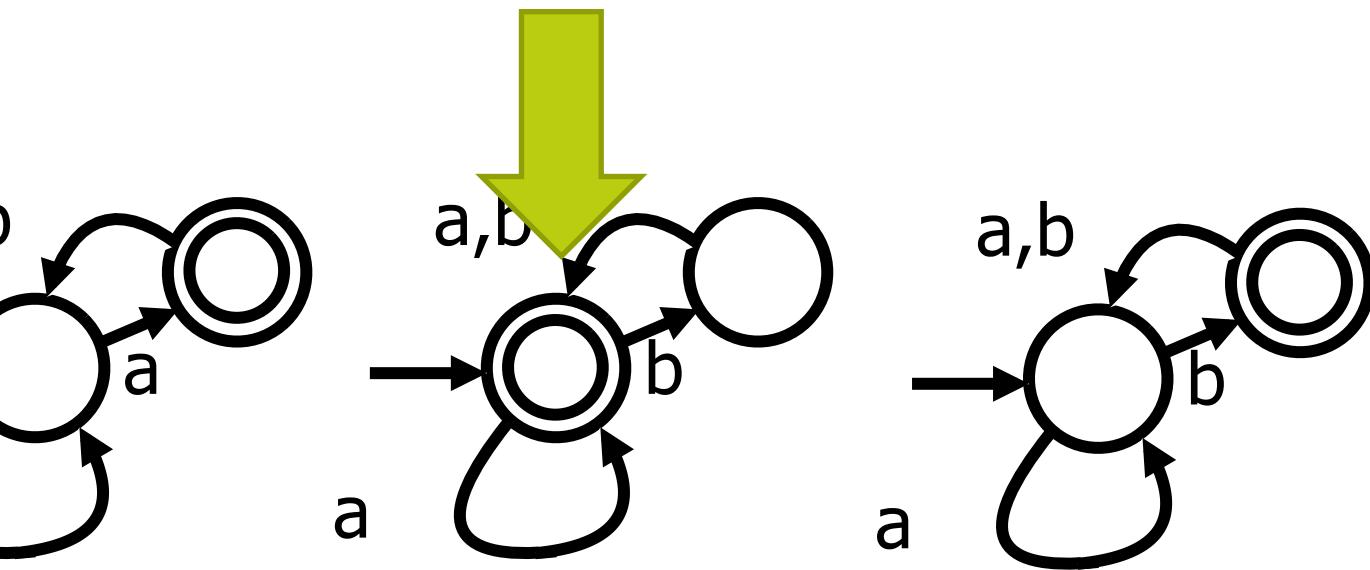
Examples: $(a,+)(aa,+)(b,-)$

Learning by enumeration

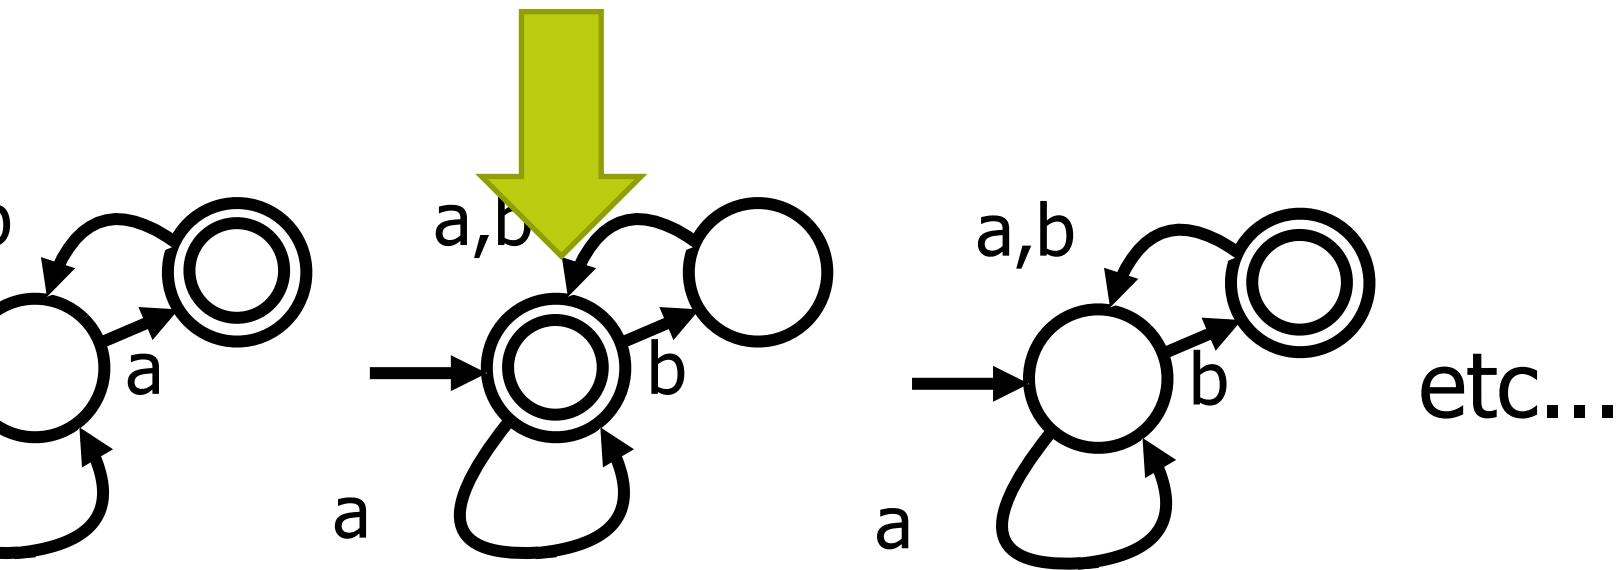


Examples: $(a,+)(aa,+)(b,-)$

Learning by enumeration



Learning by enumeration



Learning DFAs

- DFAs are learnable in the limit from labeled data
 - by enumeration using **Occam's razor**, smallest DFA first
- Are DFAs learnable in the limit from unlabeled data?
 - try learning $\{a,b\}^*$ (all strings) and all finite languages over $\{a,b\}$

Learning DFAs

- DFAs cannot be learned in the limit from unlabeled data!
 - Limit learners should at some point **converge** to C
 - $\{a,b\}^*$ contains all finite languages over $\{a,b\}$
 - Suppose a learner converges to $\{a,b\}^*$ at some point
 - At that point, the learner has seen **a finite data set D**
 - D is a finite language
 - The learner can never converge to D
- **Contradiction**

Learning DFAs from labeled data

- Input:
 - Labeled data, positive and negative traces
- Goal:
 - Find the **smallest** DFA that is **consistent** with the data
- Learnable in the limit, but NP-complete
 - proof from graph coloring
- Solutions:
 - Rely on data to be **nice**
 - Learn actively, by asking questions to a **teacher** (oracle)
 - Assume **statistical properties**, e.g., a Markov property
 - Translate the problem and use powerful **solvers**

MYHILL-NERODE AND THE HANKEL MATRIX

Myhill-Nerode

- A **necessary** and **sufficient** condition (if and only if) for a language to be regular
 - Characterises when finite state machines exist
 - We use it for learning such state machines

Myhill-Nerode

- Given a language L , and a pair of strings x and y , define a **distinguishing extension** to be a string z such that exactly one of the two strings xz and yz belong to L .
- Define a **relation R** on the strings by the rule that
 - $x R y$
 - if there is no distinguishing extension for x and y .
- It is easy to show that R is an **equivalence relation**.
- L is regular if and only if R has a **finite number of equivalence classes** (sets of R -equivalent strings).
- Moreover, the number of states of the smallest deterministic finite state automaton recognizing L is equal to the number of equivalence classes in R .

Myhill-Nerode

- Given a language L , and a pair of strings x and y , define a **distinguishing extension** to be a string z such that exactly one of the two strings xz and yz belong to L .
- Define a relation R on the strings by the rule that
 - $x R y$
 - if there is no distinguishing extension for x and y .
- It is easy to show that R is an equivalence relation.
- L is regular if and only if R has a finite number of equivalence classes (sets of R -equivalent strings).
- Moreover, the number of states of the smallest deterministic finite state automaton recognizing L is equal to the number of equivalence classes in R .

Visualization

	_	a	b	aa	ab	ba	bb	aaa	aab	aba
aa										
ba										

- $x = aa$
- $y = ba$
- $z = \text{column}$

two strings aa and ba
and possible distinguishing extensions

Visualization

$a(a|b)^*b$

	_	a	b	aa	ab	ba	bb	aaa	aab	aba
aa										
ba										

Visualization

$a(a|b)^*b$

	_	a	b	aa	ab	ba	bb	aaa	aab	aba
aa	0									
ba										

Visualization

$a(a|b)^*b$

	_	a	b	aa	ab	ba	bb	aaa	aab	aba
aa	0									
ba	0									

Visualization

$a(a|b)^*b$

	_	a	b	aa	ab	ba	bb	aaa	aab	aba
aa	0	0								
ba	0	0								

Visualization

$a(a|b)^*b$

	_	a	b	aa	ab	ba	bb	aaa	aab	aba
aa	0	0	1							
ba	0	0	0							

Visualization

$a(a|b)^*b$

	_	a	b	aa	ab	ba	bb	aaa	aab	aba
aa	0	0	1	0	1	0	1	0	1	0
ba	0	0	0	0	0	0	0	0	0	0

Myhill-Nerode

- Given a language L , and a pair of strings x and y , define a distinguishing extension to be a string z such that exactly one of the two strings xz and yz belong to L .
- Define a relation R on the strings by the rule that
 - $x R y$
 - if there is no distinguishing extension for x and y .
- It is easy to show that R is an equivalence relation.
- L is regular if and only if R has a finite number of equivalence classes (sets of R -equivalent strings).
- Moreover, the number of states of the smallest deterministic finite state automaton recognizing L is equal to the number of equivalence classes in R .

	_	a	b	aa	ab	ba	bb	aaa	aab	aba
_	0	0	0	0	1	0	0	0	1	0
a	0	0	1	0	1	0	1	0	1	0
b	0	0	0	0	0	0	0	0	0	0
aa	0	0	1	0	1	0	1	0	1	0
ab	1	0	1	0	1	0	1	0	1	0
ba	0	0	0	0	0	0	0	0	0	0
bb	0	0	0	0	0	0	0	0	0	0
aaa	0	0	1	0	1	0	1	0	1	0
aab	1	0	1	0	1	0	1	0	1	0
aba	0	0	1	0	1	0	1	0	1	0

	_	a	b	aa	ab	ba	bb	aaa	aab	aba
_	0	0	0	0	1	0	0	0	1	0
a	0	0	1	0	1	0	1	0	1	0
b	0	0	0	0	0	0	0	0	0	0
aa	0	Hankel matrix of language L								1 0
ab	1									1 0
ba	0									0 0
bb	0	0	0	0	0	0	0	0	0	0
aaa	0	0	1	0	1	0	1	0	1	0
aab	1	0	1	0	1	0	1	0	1	0
aba	0	0	1	0	1	0	1	0	1	0

Hankel matrix

- Rows are prefixes
- Columns are suffixes
- Row-column (x,y) is:
 - 1 if xy in L
 - 0 if xy not in L

	_	a	b	aa	ab	ba	bb	aaa	aab	aba
_	0	0	0	0	1	0	0	0	1	0
a	0	0	1	0	1	0	1	0	1	0
b	0	0	0	0	0	0	0	0	0	0
aa	0	0	1	0	1	0	1	0	1	0
ab	1	0	1	0	1	0	1	0	1	0
ba	0	0	0	0	0	0	0	0	0	0
bb	0	0	0	0	0	0	0	0	0	0
aaa	0	0	1	0	1	0	1	0	1	0
aab	1	0	1	0	1	0	1	0	1	0
aba	0	0	1	0	1	0	1	0	1	0

	_	a	b	aa	ab	ba	bb	aaa	aab	aba
_	0	0	0	0	1	0	0	0	1	0
a	0	0	1	0	1	0	1	0	1	0
b	0	0	0	0	0	0	0	0	0	0
aa	0	0	1	0	1	0	1	0	1	0
ab	1	0	1	0	1	0	1	0	1	0
ba	0	0	0	0	0	0	0	0	0	0
bb	0	0	0	0	0	0	0	0	0	0
aaa	0	0	1	0	1	0	1	0	1	0
aab	1	0	1	0	1	0	1	0	1	0
aba	0	0	1	0	1	0	1	0	1	0

	_	a	b	aa	ab	ba	bb	aaa	aab	aba
_	0	0	0	0	1	0	0	0	1	0
a	0	0	1	0	1	0	1	0	1	0
b	0	0	0	0	0	0	0	0	0	0
aa	0	0	1	0	1	0	1	0	1	0
ab	1	0	1	0	1	0	1	0	1	0
ba	0	0	0	0	0	0	0	0	0	0
bb	0	0	0	0	0	0	0	0	0	0
aaa	0	0	1	0	1	0	1	0	1	0
aab	1	0	1	0	1	0	1	0	1	0
aba	0	0	1	0	1	0	1	0	1	0

	_	a	b	aa	ab	ba	bb	aaa	aab	aba
_	0	0	0	0	1	0	0	0	1	0
a	0	0	1	0	1	0	1	0	1	0
b	0	0	0	0	0	0	0	0	0	0
aa	0	0	1	0	1	0	1	0	1	0
ab	1	0	1	0	1	0	1	0	1	0
ba	0	0	0	0	0	0	0	0	0	0
bb	0	0	0	0	0	0	0	0	0	0
aaa	0	0	1	0	1	0	1	0	1	0
aab	1	0	1	0	1	0	1	0	1	0
aba	0	0	1	0	1	0	1	0	1	0

	_	a	b	aa	ab	ba	bb	aaa	aab	aba
_	0	0	0	0	1	0	0	0	1	0
a	0	0	1	0	1	0	1	0	1	0
b	0	0	0	0	0	0	0	0	0	0
aa	0	0	1	0	1	0	1	0	1	0
ab	1	0	1	0	1	0	1	0	1	0
ba	0	0	0	0	0	0	0	0	0	0
bb	0	0	0	0	0	0	0	0	0	0
aaa	0	0	1	0	1	0	1	0	1	0
aab	1	0	1	0	1	0	1	0	1	0
aba	0	0	1	0	1	0	1	0	1	0

	_	a	b	aa	ab	ba	bb	aaa	aab	aba
_	0	0	0	0	1	0	0	0	1	0
a	0	0	1	0	1	0	1	0	1	0
b	0	0	0	0	0	0	0	0	0	0
aa	x R y if row x and row y have the same color									0
ab										0
ba										0
bb	0	0	0	0	0	0	0	0	0	0
aaa	0	0	1	0	1	0	1	0	1	0
aab	1	0	1	0	1	0	1	0	1	0
aba	0	0	1	0	1	0	1	0	1	0

Myhill-Nerode

- Given a language L , and a pair of strings x and y , define a distinguishing extension to be a string z such that exactly one of the two strings xz and yz belong to L .
- Define a relation R on the strings by the rule that
 - $x R y$
 - if there is no distinguishing extension for x and y .
- **It is easy to show that R is an equivalence relation.**
- L is regular if and only if R has a finite number of equivalence classes (sets of R -equivalent strings).
- Moreover, the number of states of the smallest deterministic finite state automaton recognizing L is equal to the number of equivalence classes in R .

Myhill-Nerode

- Given a language L , and a pair of strings x and y , define a distinguishing extension to be a string z such that exactly one of the two strings xz and yz belong to L .
- Define a relation R on the strings by the rule that
 - $x R y$
 - if there is no distinguishing extension for x and y .
- It is easy to show that R is an equivalence relation.
- L is regular if and only if R has a **finite number of equivalence classes** (sets of R -equivalent strings).
- Moreover, the number of states of the smallest deterministic finite state automaton recognizing L is equal to the number of equivalence classes in R .

	_	a	b	aa	ab	ba	bb	aaa	aab	aba
_	0	0	0	0	1	0	0	0	1	0
a	0	0	1	0	1	0	1	0	1	0
b	0	0	0	0	0	0	0	0	0	0
aa	0	0	1	0	1	0	1	0	1	0
ab	1	0	1	0	1	0	1	0	1	0
ba	0	0	0	0	0	0	0	0	0	0
bb	0	0	0	0	0	0	0	0	0	0
aaa	0	0	1	0	1	0	1	0	1	0
aab	1	0	1	0	1	0	1	0	1	0
aba	0	0	1	0	1	0	1	0	1	0

	_	a	b	aa	ab	ba	bb	aaa	aab	aba
_	0	0	0	0	1	0	0	0	1	0
a	0	0	1	0	1	0	1	0	1	0
b	0	1	0	1	0	1	0	1	0	1
aa	0	0	1	0	1	0	1	0	1	0
aab	1	0	1	0	1	0	1	0	1	0
aba	0	0	1	0	1	0	1	0	1	0

A language is regular if and only if the Hankel matrix contains a **finite number of distinct rows**, i.e., colors

Myhill-Nerode

- Given a language L , and a pair of strings x and y , define a distinguishing extension to be a string z such that exactly one of the two strings xz and yz belong to L .
- Define a relation R on the strings by the rule that
 - $x R y$
 - if there is no distinguishing extension for x and y .
- It is easy to show that R is an equivalence relation.
- L is regular if and only if R has a finite number of equivalence classes (sets of R -equivalent strings).
- Moreover, the number of states of the **smallest deterministic finite state automaton** recognizing L is equal to the number of equivalence classes in R .

	_	a	b	aa	ab	ba	bb	aaa	aab	aba
_	0	0	0	0	1	0	0	0	1	0
a	0	0	1	0	1	0	1	0	1	0
b	0	1	0	1	0	1	0	1	0	1
aa	0	0	1	0	1	0	1	0	1	0
ab	1	0	1	0	1	0	1	0	1	0
ba	0	0	1	0	1	0	1	0	1	0
bb	0	1	0	1	0	1	0	1	0	1
aaa	0	0	1	0	1	0	1	0	1	0
aab	1	0	1	0	1	0	1	0	1	0
aba	0	0	1	0	1	0	1	0	1	0

The number of state of the smallest DFA for L is
the number of colors in the Hankel matrix

Constructing a state machine

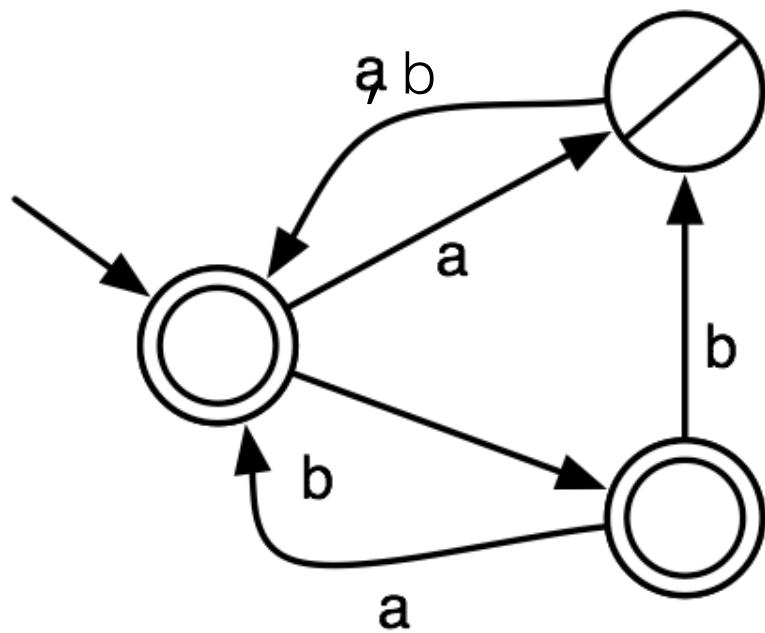
- Given a Hankel matrix for L
- What is the smallest state machine (DFA) for L ?

Q: what does this state machine look like?

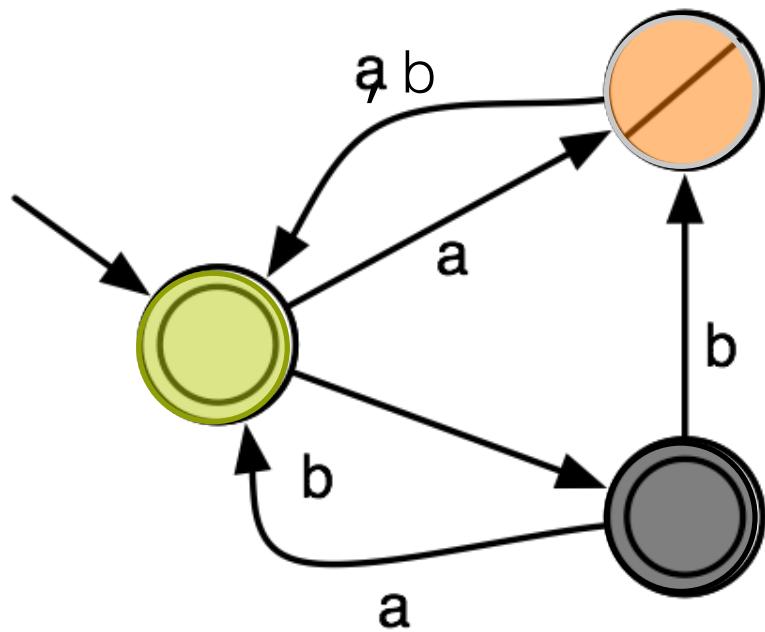
	_	a	b	aa	ab	ba	bb	aaa	aab	aba
_	1	0	1	1	1	1	0	0	1	0
a	0	1	1	0	1	0	1	1	1	1
b	1	1	0	0	1	1	1	1	1	1
aa	1	0	1	1	1	1	0	0	1	0
ab	1	0	1	1	1	1	0	0	1	0
ba	1	0	1	1	1	1	0	0	1	0
bb	0	1	1	0	1	0	1	0	1	1
aaa	0	1	1	0	1	0	1	1	1	1
aab	1	1	0	0	1	1	1	1	1	1
aba	0	1	1	0	1	0	1	1	1	1

	_	a	b	aa	ab	ba	bb	aaa	aab	aba
_	1	0	1	1	1	1	0	0	1	0
a	0	1	1	0	1	0	1	1	1	1
b	1	1	0	0	1	1	1	1	1	1
aa	1	0	1	1	1	1	0	0	1	0
ab	1	0	1	1	1	1	0	0	1	0
ba	1	0	1	1	1	1	0	0	1	0
bb	0	1	1	0	1	0	1	0	1	1
aaa	0	1	1	0	1	0	1	1	1	1
aab	1	1	0	0	1	1	1	1	1	1
aba	0	1	1	0	1	0	1	1	1	1

Answer



Answer



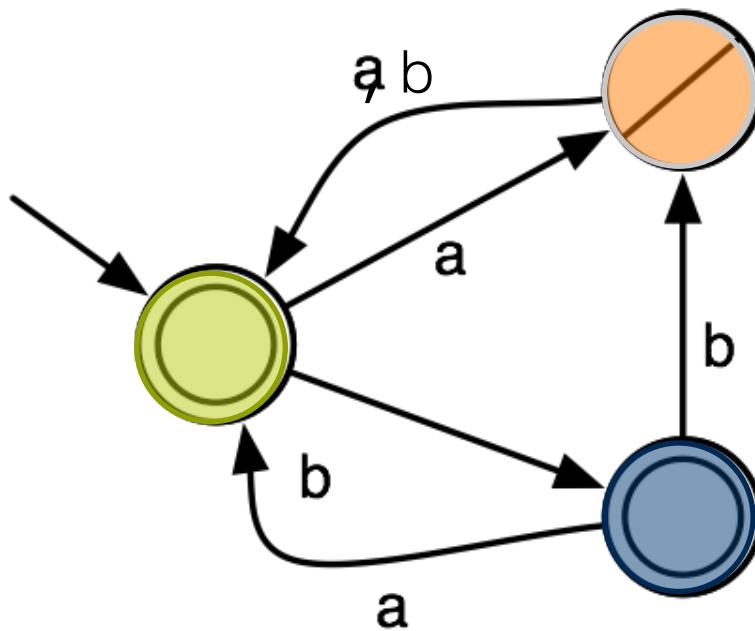
	_	a	b	aa	ab	ba	bb	aaa	aab	aba
_	1	0	1	1	1	1	0	0	1	0
a	0	1	1	0	1	0	1	1	1	1
b	1	1	0	0	1	1	1	1	1	1
aa	1	0	1	1	1	1	0	0	1	0
ab	1	0	1	1	1	1	0	0	1	0
ba	1	0	1	1	1	1	0	0	1	0
bb	0	1	1	0	1	0	1	0	1	1
aaa	0	1	1	0	1	0	1	1	1	1
aab	1	1	0	0	1	1	1	1	1	1
aba	0	1	1	0	1	0	1	1	1	1

Look at colors of Hankel Matrix rows

—
a
b
aa
ab
ba
bb
aaa
aab
aba

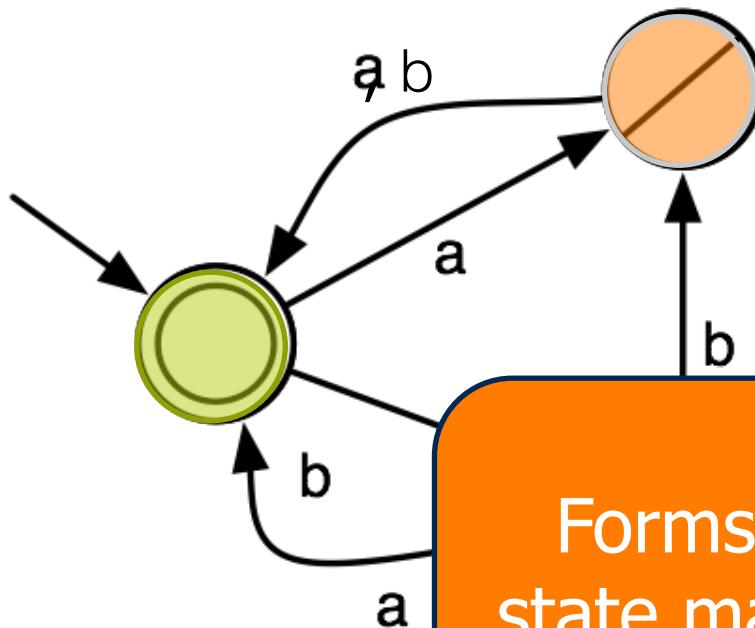
–
a
b
aa
ab
ba
bb

access strings and *one-letter extensions*
provide sufficient information



—
a
b
aa
ab
ba
bb

access strings and *one-letter extensions*
provide sufficient information



Forms the basis for
state machine learning
algorithms

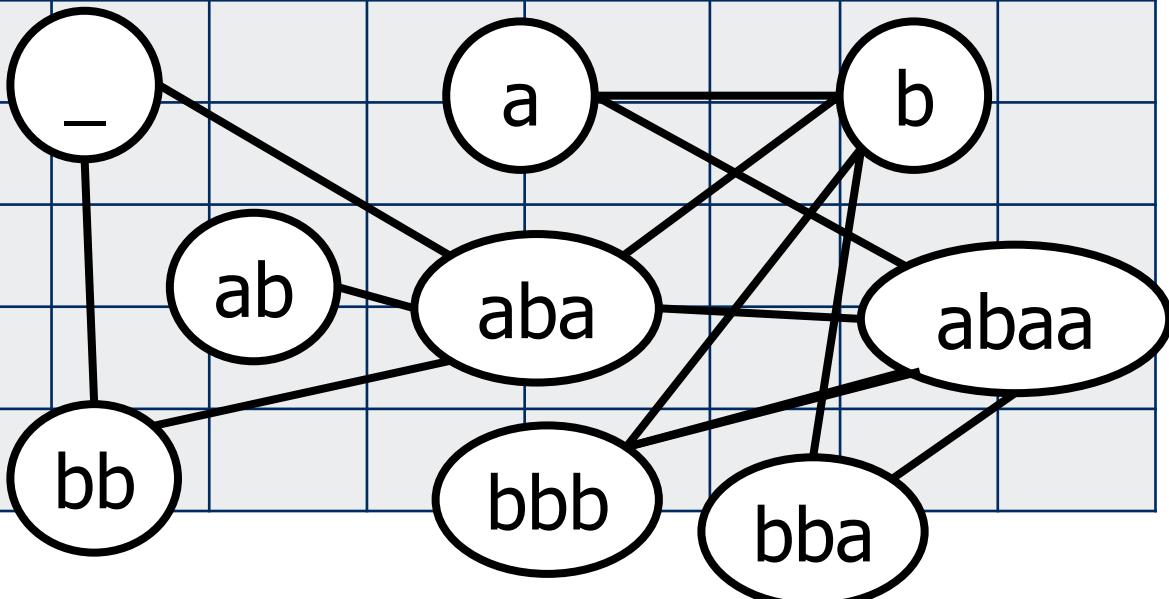
STATE MACHINE LEARNING

(a,+) (b,-) (aba,+) (abaa,-) (bbb,+) (bba,+)

	<u>_</u>	a	b	ba	aba	aa	baa	abaa	bb	bbb	bba
<u>_</u>		1	0		1			0		1	1
a	1			1			0				
b	0			1					1		
ab		1				0					
aba	1	0									
abaa	0										
bb		1	1								
bbb	1										
bba	1										

(a,+) (b,-) (aba,+) (abaa,-) (bbb,+) (bba,+)

	<u>_</u>	a	b	ba	aba	aa	baa	abaa	bb	bbb	bba
<u>_</u>	1	0		1				0		1	1
a	1			1			0				
b	0			1					1		
ab		1				0					
aba	1	0									
abaa	0										
bb		1	1								
bbb	1										
bba	1										



(a,+) (b,-) (aba,+) (abaa,-) (bbb,+) (bba,+)

	<u>_</u>	a	b	ba	aba	aa	baa	abaa	bb	bbb	bba
<u>_</u>	1	0		1				0		1	1
a	1			1			0				
b	0			1					1		
ab		1				0					
aba	1	0									
abaa	0										
bb		1	1								
bbb	1										
bba	1										

(a,+) (b,-) (aba,+) (abaa,-) (bbb,+) (bba,+)

a b ba aba aa

baa abaa

bb bbb bba

0 1 1

0 1

ab 1

0

aba 1 0

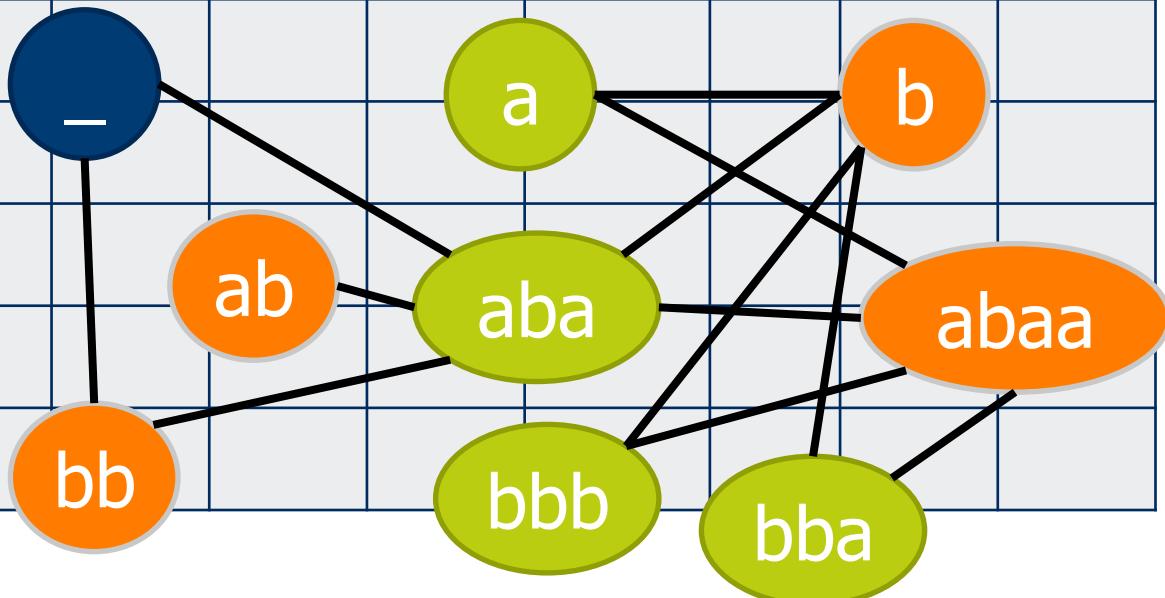
abaa 0

bb 1 1

bbb 1

bba 1

Try to make state machine...



(a,+) (b,-) (aba,+) (abaa,-) (bbb,+) (bba,+)

Try to make state machine...

It is **inconsistent!**

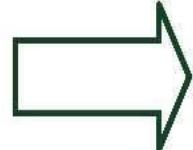
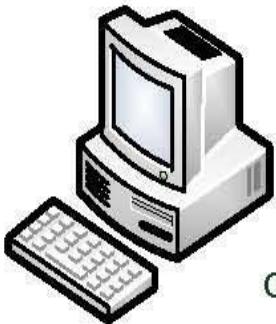
(b and bb end in same state,
but bbb does not)

a	b	ba	aba	aa	baa	abaa	bb	bbb	bba
					0			1	1
					0			1	
ab		1			0				
aba	1	0							
abaa	0								
bb		1	1						
bbb	1								
bba	1								

A state transition diagram is shown below the table. It consists of seven states represented by circles: a blue circle labeled '-' at the top left, and six colored circles (orange and green) labeled 'ab', 'aba', 'aa', 'bb', 'bba', and 'b'. Transitions are indicated by arrows between states. The '-' state has arrows pointing to 'ab' (orange), 'aba' (green), and 'aa' (green). The 'ab' state has arrows pointing to 'aba' (green), 'aa' (green), and 'bb' (orange). The 'aa' state has an arrow pointing to 'bba' (green). The 'bb' state has arrows pointing to 'bba' (green) and 'b' (orange). The 'bba' state has an arrow pointing to 'b' (orange). The 'b' state has an arrow pointing back to '-'. The 'aba' state has no outgoing arrows.

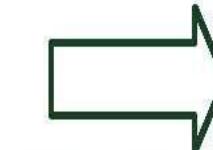
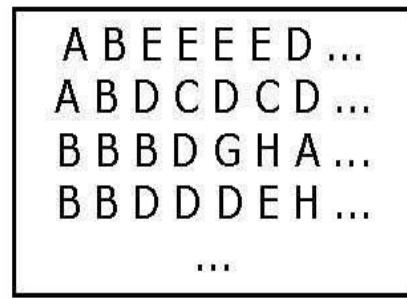
We need more information...

software system

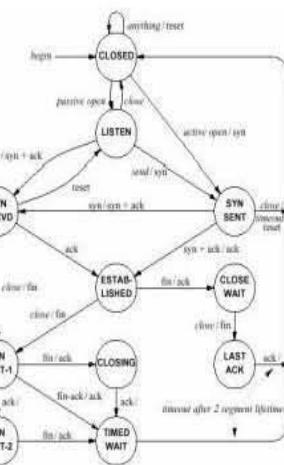


system call or
communication logs

execution traces



state machine
learning



software
model

Queries (input)



Active learning

	<u>_</u>	a	b	ba	aba	aa	baa	abaa	bb	bbb	bba
<u>_</u>	1	0		1				0		1	1
a	1			1			0				
b	0			1					1		
ab		1				0					
aba	1	0									
abaa	0										
bb		1	1								
bbb	1										
bba	1										

What to ask?

Ask: bb

	<u>_</u>	a	b	ba	aba	aa	baa	abaa	bb	bbb	bba
<u>_</u>	1	0		1				0		1	1
a	1			1			0				
b	0			1					1		
ab		1				0					
aba	1	0									
abaa	0										
bb		1	1								
bbb	1										
bba	1										

Ask: bb

	_	a	b	ba	aba	aa	baa	abaa	bb	bbb	bba
_	1	0		1				0		1	1
a	1			1			0				
b	0		1	1					1		
ab		1				0					
aba	1	0									
abaa	0										
bb	1	1	1								
bbb	1										
bba	1										

If + _ distinguishes between b and bb

Ask: bb

	-	a	b	ba	aba	aa	baa	abaa	bb	bbb	bba
-	1	0		1				0		1	1
a	1			1			0				
b	0		0	1					1		
ab		1				0					
aba	1	0									
abaa	0										
bb	0	1	1								
bbb	1										
bba	1										

If - b distinguishes between b and bb

Active learning with L*

- Aims to discover:
 - **access strings**, smallest strings for every color
 - **one-letter extensions**, and
 - their **distinguishing suffixes**
- Maintains a distinguishing table
- Which distinguishing suffixes is determined during learning

	—	a
—	1	0
a	0	1
b	1	1
aa	1	0
ab	1	0
ba	1	0
bb	0	1

Active learning with L*

- Start with an empty table, with one access string and one distinguishing extension:
 - _ and _
- Ask membership queries for _ and its one letter extensions:
 - _ in L? \rightarrow yes
 - a in L? \rightarrow no
 - b in L? \rightarrow yes

	-
-	1
a	0
b	1

Active learning with L^*

- Can we construct a state machine?
 - No!
 - a points to a new state
- Ask membership queries for its one-letter extensions:
 - aa in L ? \rightarrow yes
 - ab in L ? \rightarrow yes

—	—
—	1
a	0
b	1
aa	1
ab	1

Active learning with L^*

- Can we construct a state machine?
 - Yes!
- Ask an equivalence query:
 - $L(M) = L? \rightarrow$ no, M for instance accepts `aabbaa`

—	—
—	1
a	0
b	1
aa	1
ab	1

Active learning with L^*

- $C = aabbaa, -$
- Where does M make a mistake?
 - Replace c prefixes with reached access strings
- Ask membership queries:
 - $_bbaa$ in L ? \rightarrow no
 - $_baa$ in L ? \rightarrow no
 - $_aa$ in L ? \rightarrow yes

	-
	1
-	1
a	0
b	1
aa	1
ab	1

Active learning with L^*

- $C = aabbaa, -$
- Where does M make a mistake?
 - Replace c prefixes with reached access strings
- Ask membership queries:
 - $_bbaa$ in L ? \rightarrow no
 - $_baa$ in L ? \rightarrow no
 - $_aa$ in L ? \rightarrow yes

	-
-	1
a	0
b	1
aa	1
ab	1

aa is a new
distinguishing extension!

Active learning with L*

- Add the column aa
- Ask membership queries to complete the table

	–	aa
–	1	1
a	0	0
b	1	0
aa	1	1
ab	1	1

Active learning with L*

- Add the column aa
- Ask membership queries to complete the table
- b points to a new state
- Add its one-letter extensions, ask membership queries to complete

	—	aa
—	1	1
a	0	0
b	1	0
aa	1	1
ab	1	1
ba	1	1
bb	0	0

Active learning with L*

- Add the column aa
- Ask membership queries

The table defines a state machine, ask an equivalence query → yes!

Done

complete

	—	aa
—	1	1
a	0	0
b	1	0
aa	1	1
ab	1	1
ba	1	1
bb	0	0

Active learning with L*

1. Initialize the distinguishing table T
2. Complete T until it defines a state machine M
3. While(Eq(M) gives a counterexample c)
 1. Use Mem() and c to find a new distinguishing extension d
 2. Add d to T
 3. Complete T until it defines a state machine M

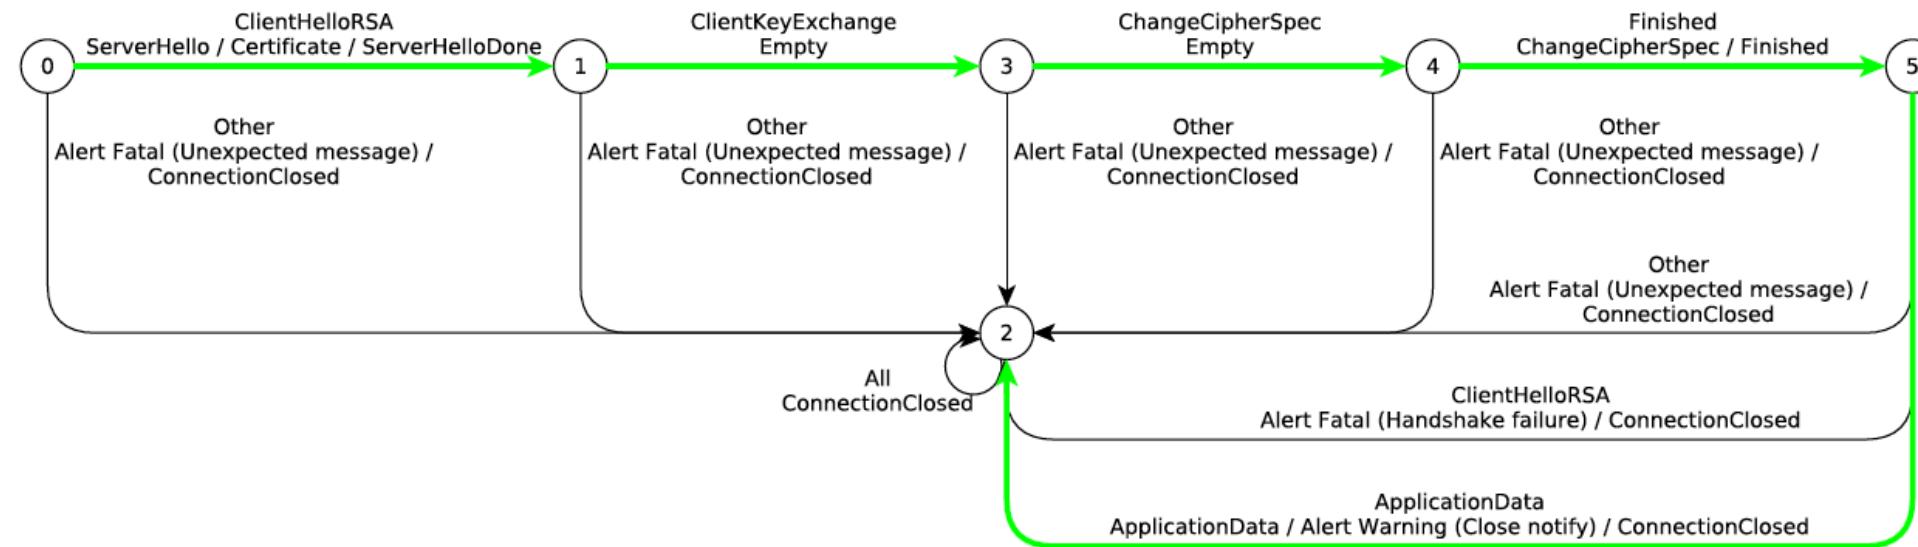
L^* - some theory

- L^* always finds a state machine for the target language
 - A distinguishing extension can always be found using the counterexample
 - Completing T terminates since we only have to add one-letter extensions once for every color
 - Once all states are identified, $\text{Eq}(M)$ will return yes
- L^* runs in polynomial time and from polynomial data!
 - Both $O(n^2)$, n number of states of smallest target

Active automaton learning

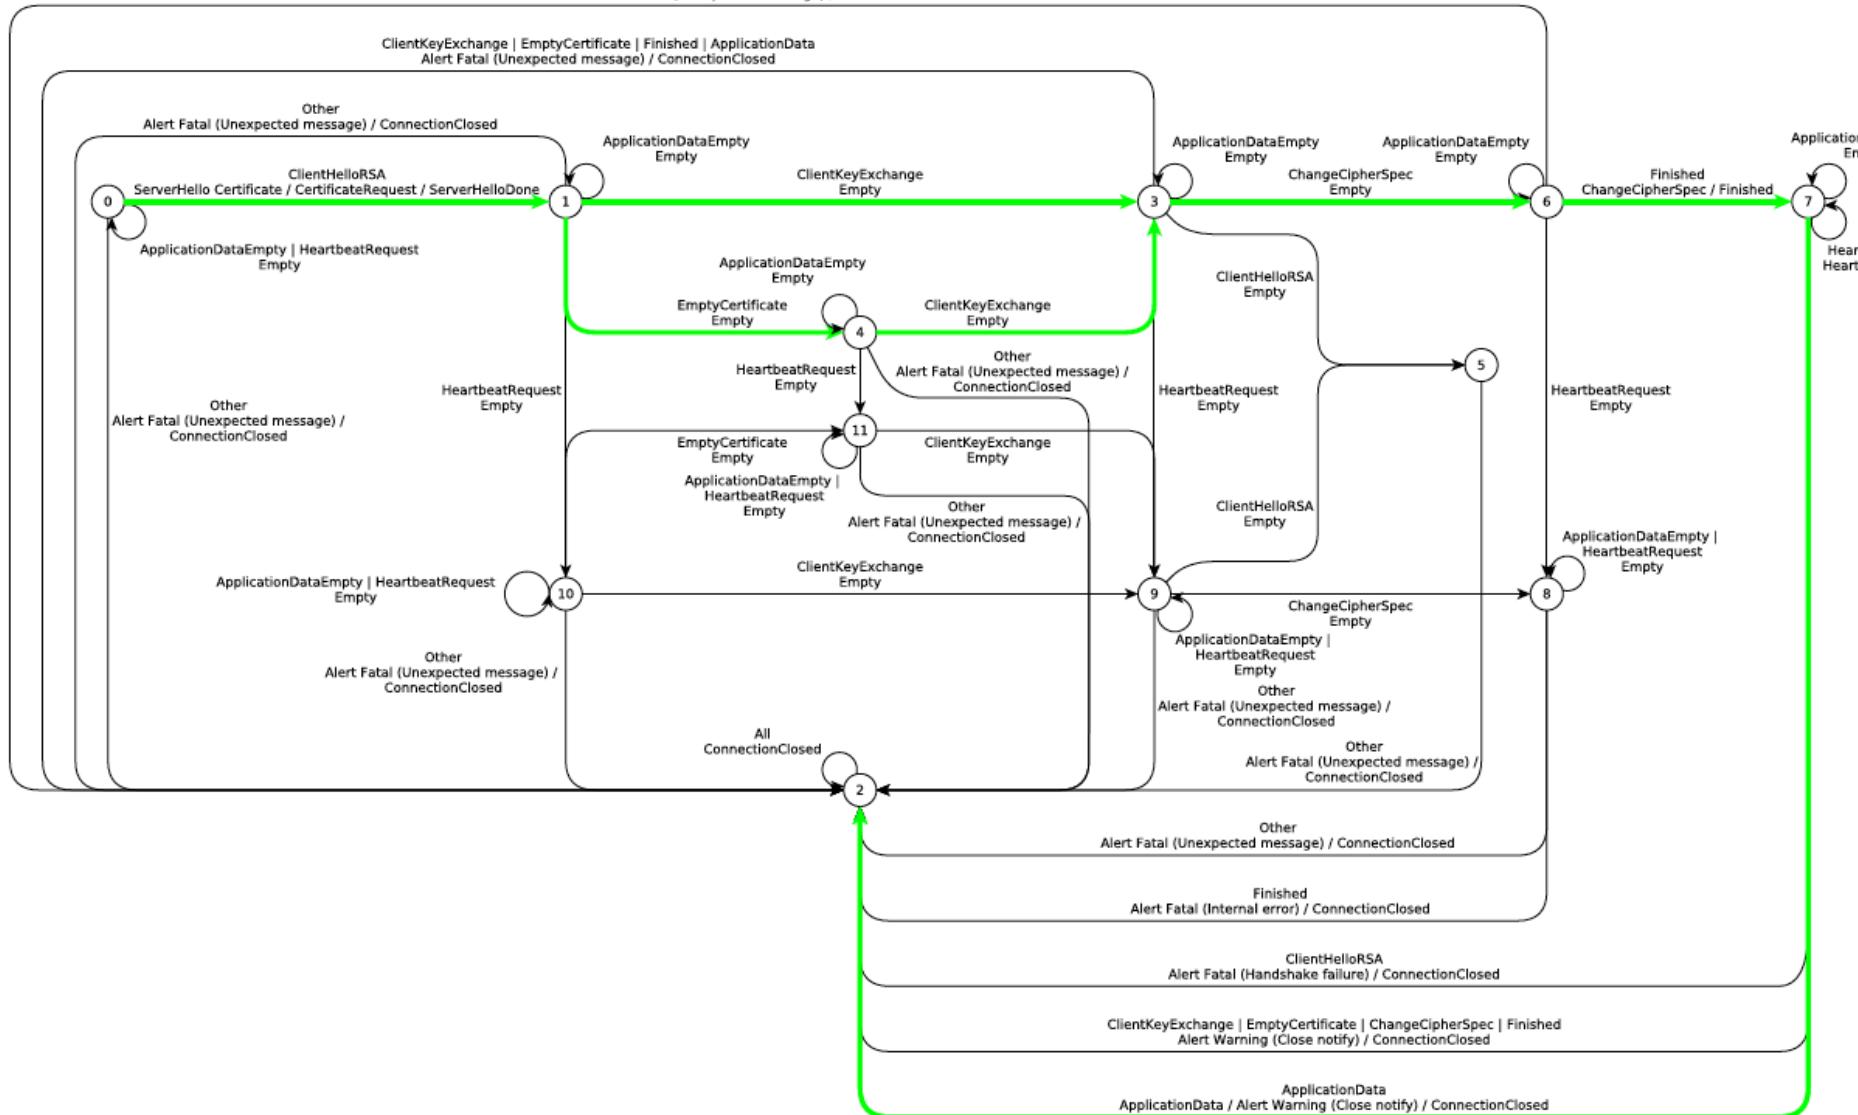
- Learnlib is an L* implementation that can be applied to software
 - Uses model-based testing methods for answering equivalence queries
 - We are currently combining this with fuzzing for security purposes
- The learned models allow to:
 - Investigate the model for any implementation flaws
 - Compare models from different implementations
 - ...
- Demo...

TLS RSA BSAFE

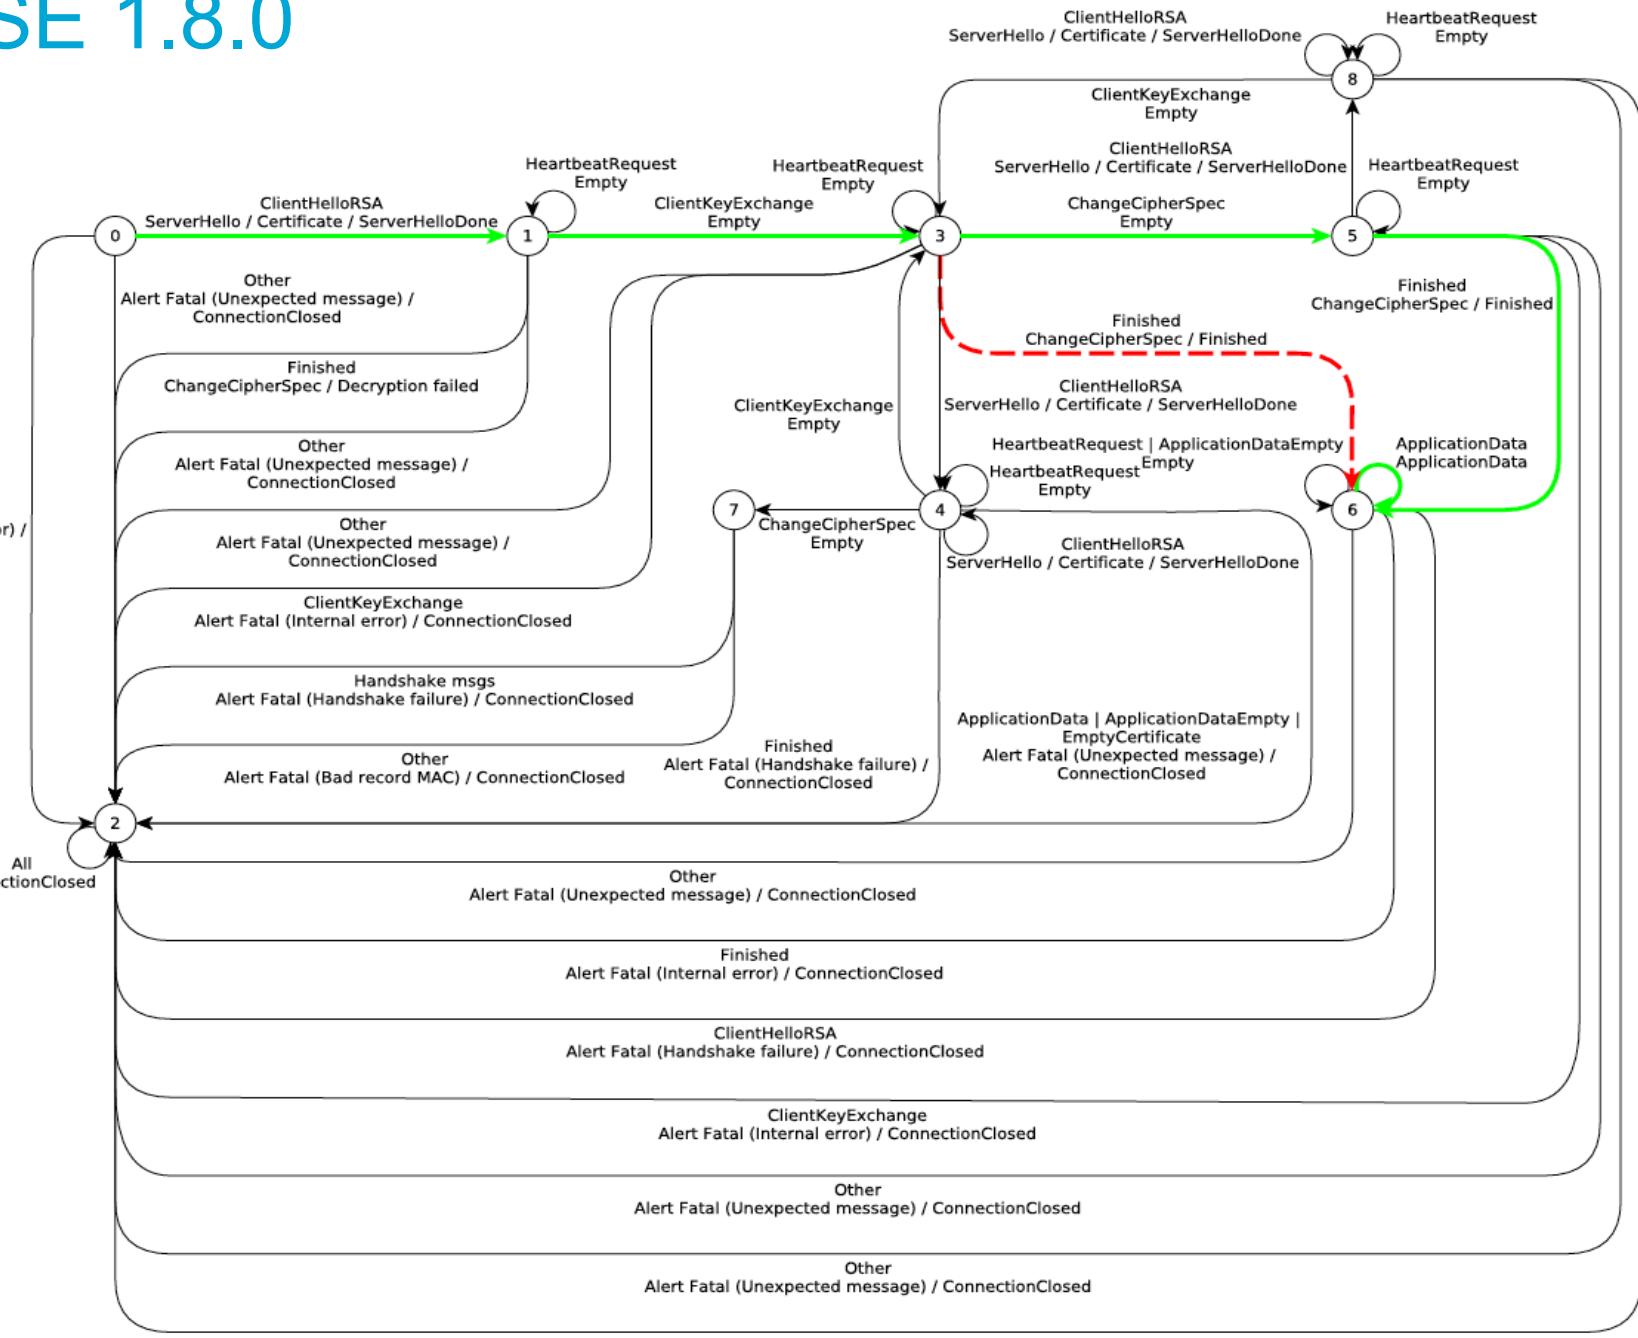


GNU TLS 3.3.8

`ClientHelloRSA | ClientKeyExchange | EmptyCertificate | ChangeCipherSpec | ApplicationData
Alert Fatal (Unexpected message) / ConnectionClosed`



JSSE 1.8.0

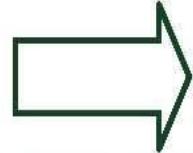
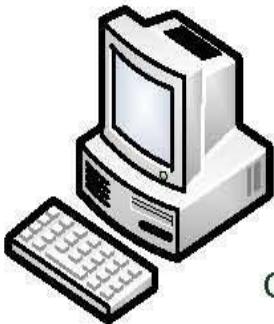


Finished Alert Fatal (Handshake failure) / ConnectionClosed

Joeri de Ruiter & Erik Poll 2011

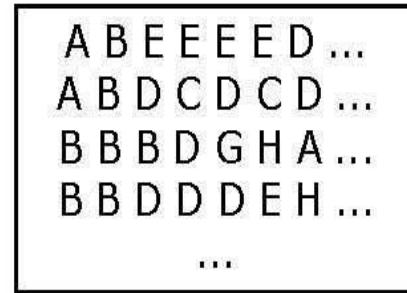
Use the information we have...

software system

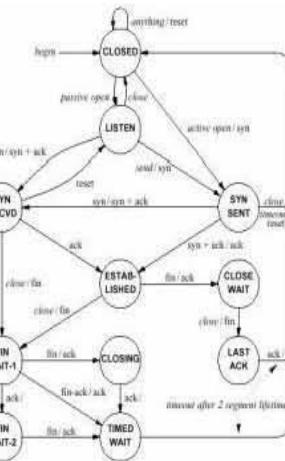


system call or
communication logs

execution traces



state machine
learning

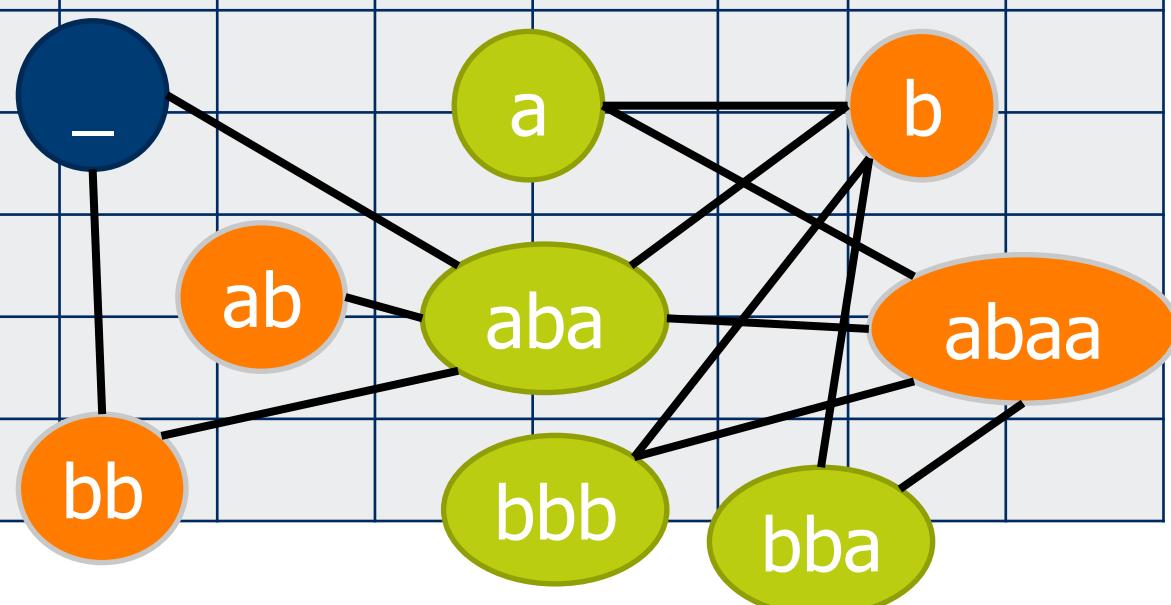


software
model

Passive learning

Key issue: coloring leads to inconsistencies

	-	a	b	ba	aba	aa	baa	abaa	bb	bbb	bba
-	1	0			1			0		1	1
a	1				1			0			
b	0				1					1	
ab		1					0				
aba	1	0									
abaa	0										
bb		1	1								
bbb	1										
bba	1										



Key issue: coloring leads to inconsistencies

	—	a	b	ba	aba	aa	baa	abaa	bb	bbb	bba
—	1	0			1			0		1	1
ab		1					0				1
aba	1	0									
abaa	0										
bb		1	1								
bbb	1										
bba	1										

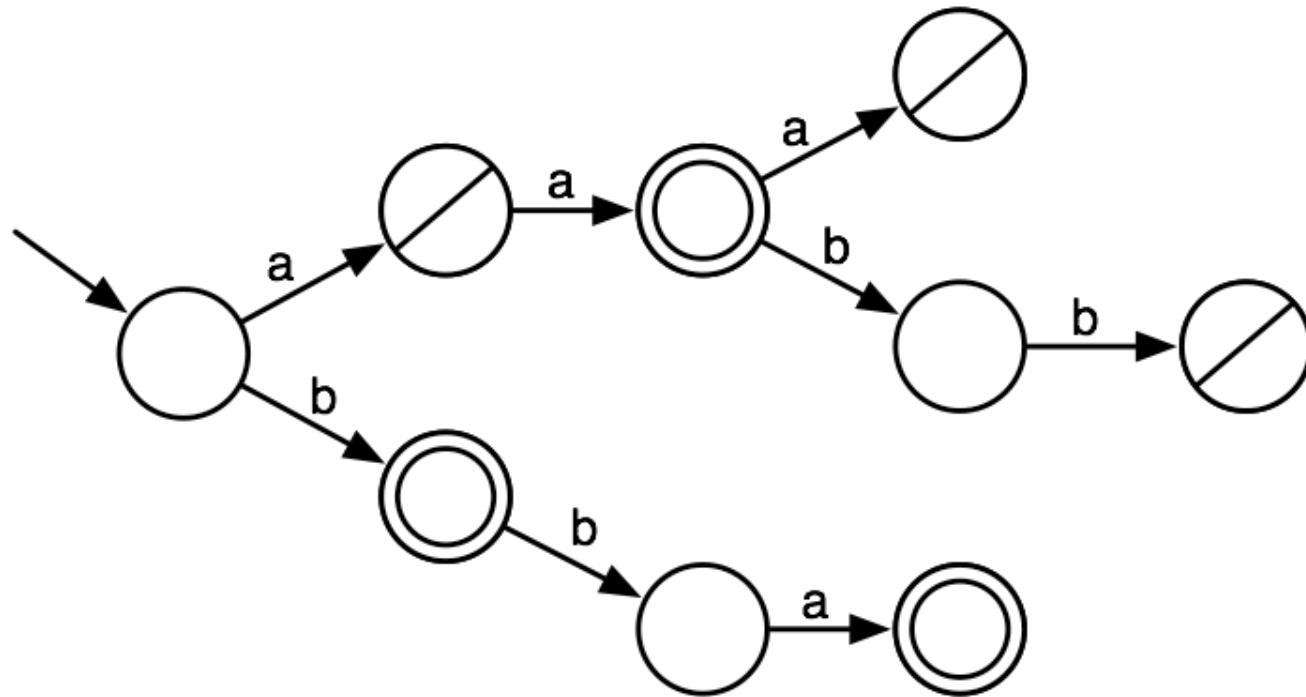
Oh, and it is NP-hard...

```
graph TD; -(( )) -- 1 --> ab((ab)); ab -- 1 --> aba((aba)); aba -- 1 --> baa((baa)); aba -- 1 --> abaa((abaa)); aba -- 1 --> bb((bb)); aba -- 1 --> bbb((bbb)); aba -- 1 --> bba((bba)); baa -- 1 --> abaa; baa -- 1 --> bb; baa -- 1 --> bbb; baa -- 1 --> bba; baa -- 1 --> -; abaa -- 1 --> ab; abaa -- 1 --> b; abaa -- 1 --> bb; abaa -- 1 --> bbb; abaa -- 1 --> bba;
```

Solution: state merging

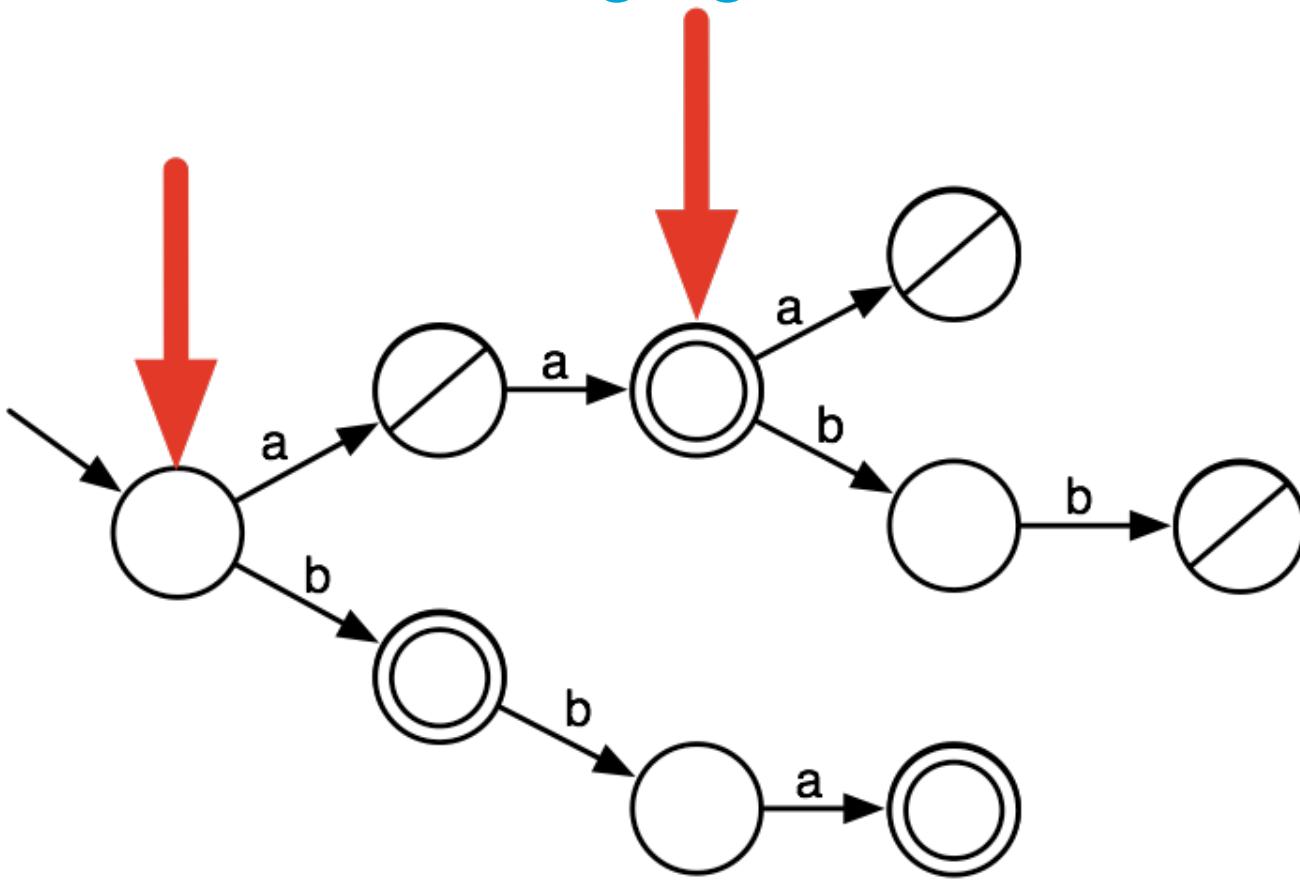
- Very effective **greedy** method
 - Initially assume a large DFA
 - Iteratively combine similar states
 - Until no more state can be combined
- State-of-the-art since 1998

State merging



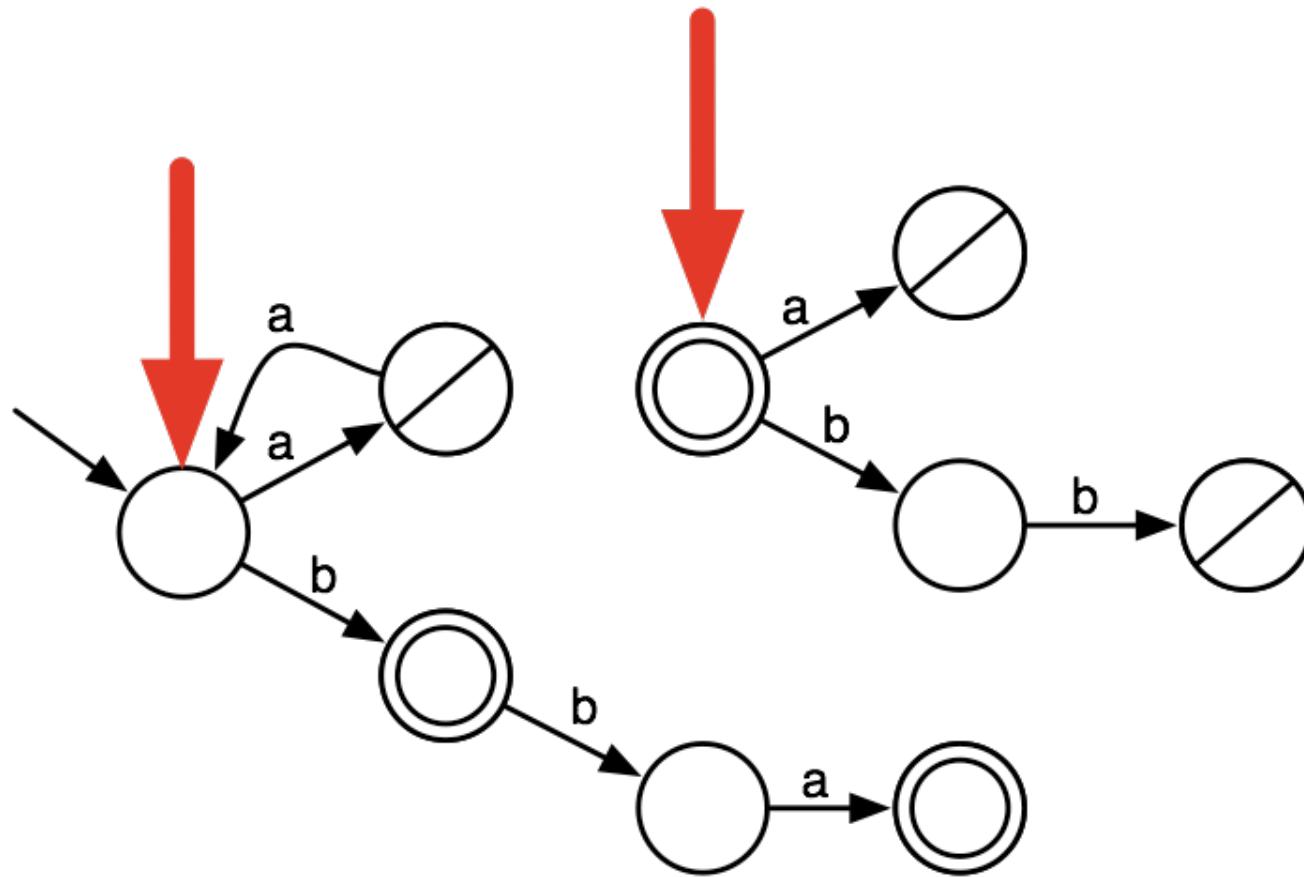
positive data: aa, b, bba; negative data: a, aaa, aabb
represented as a **prefix tree**

State merging



State merging:
select two nodes

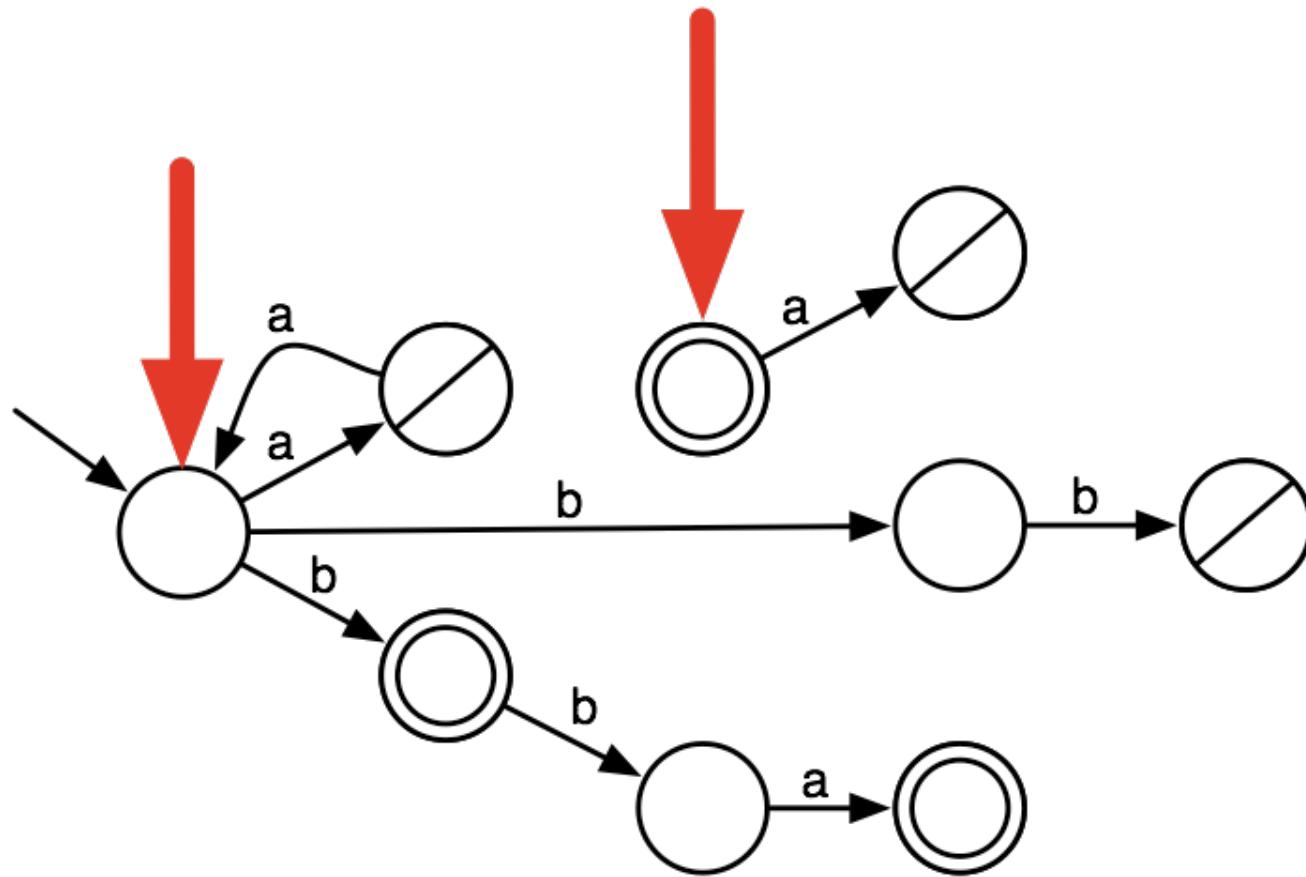
State merging



State merging:

move **input** transitions from one state to the other

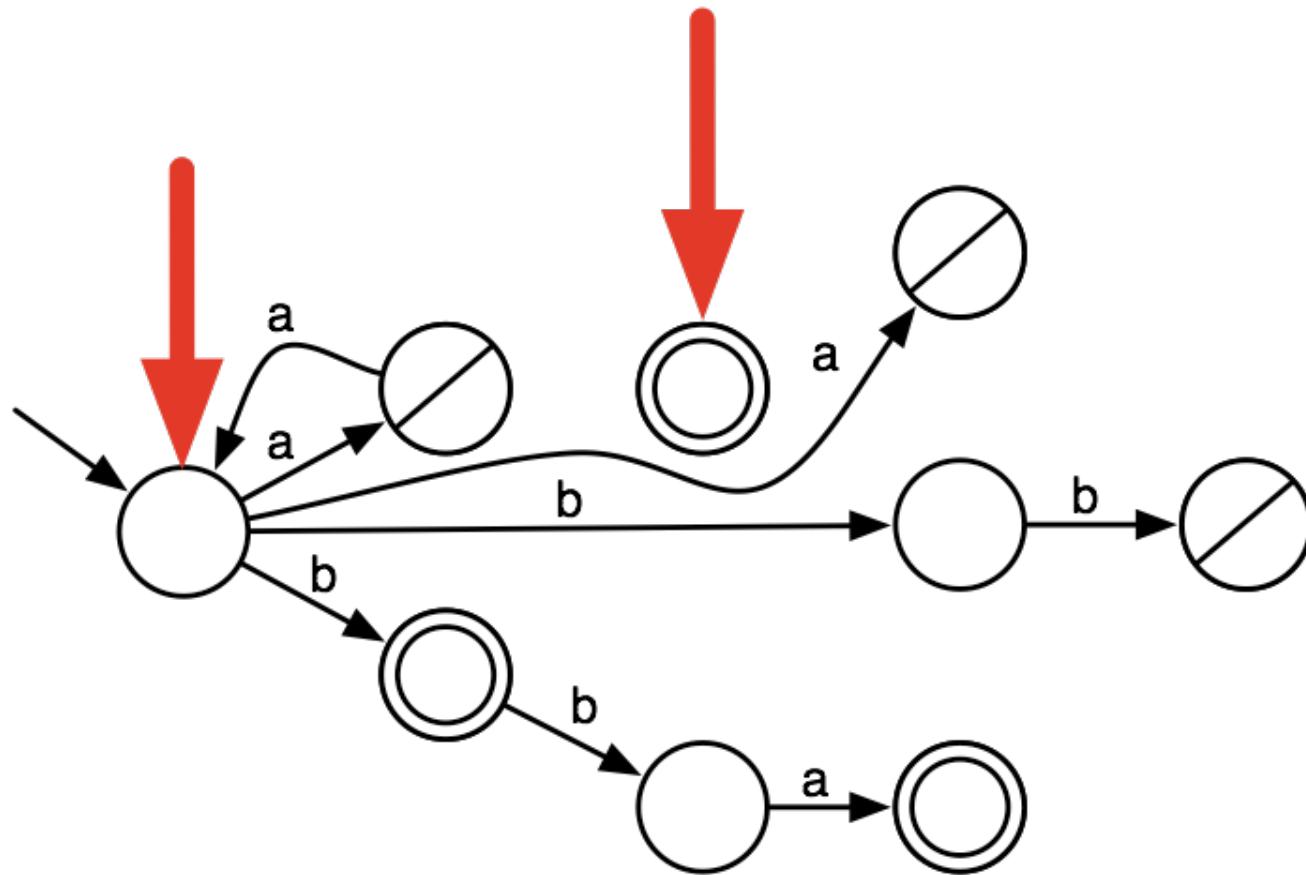
State merging



State merging:

move **output** transitions from one state to the other

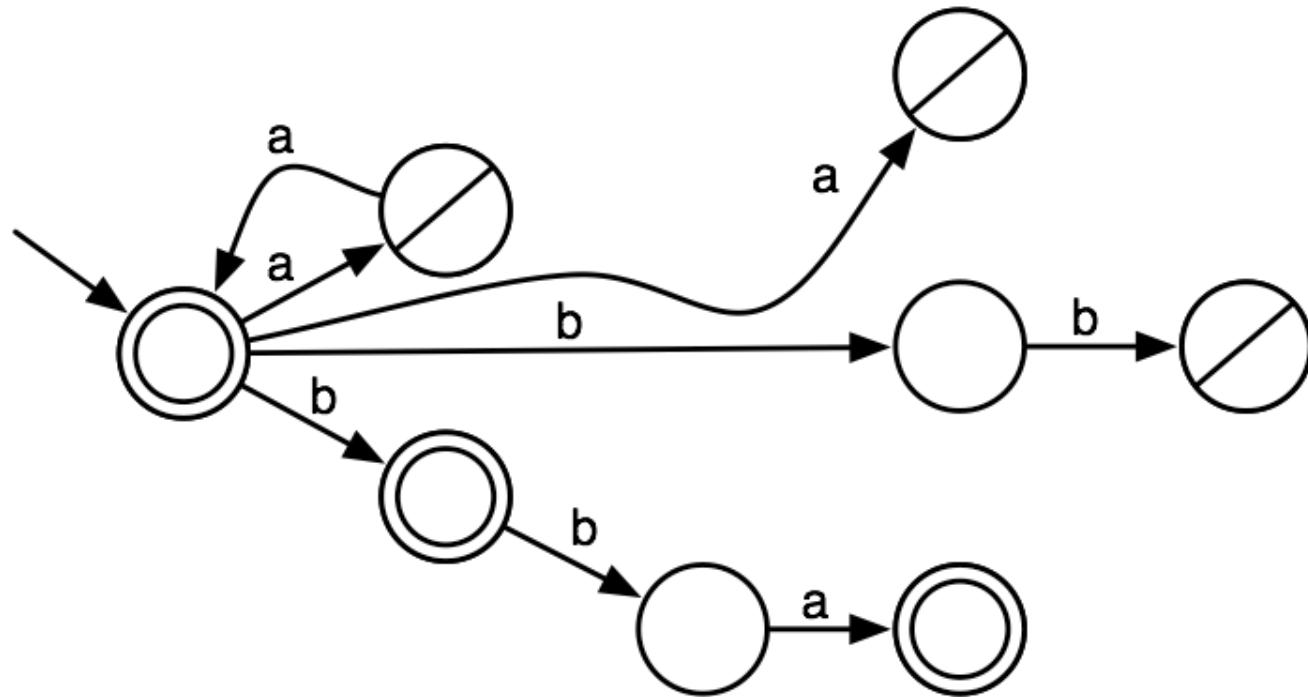
State merging



State merging:

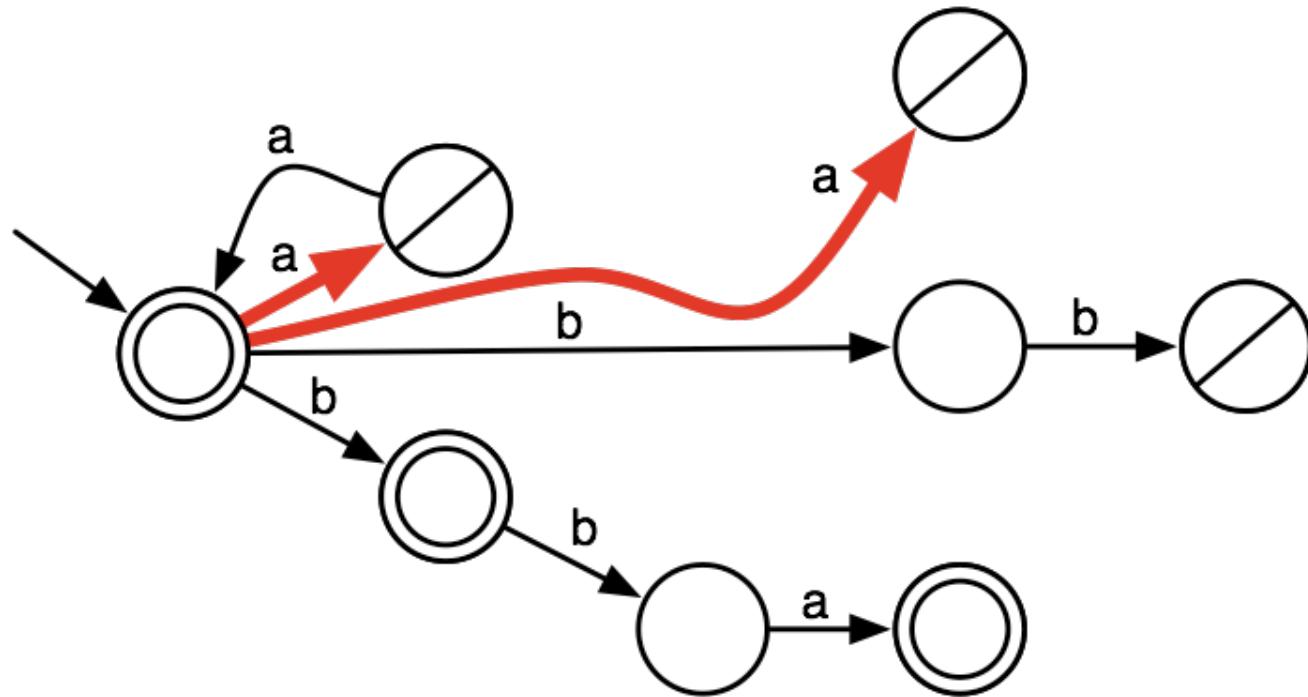
move **output** transitions from one state to the other

State merging



State merging:
delete the obsolete state, **maintain** positive/negative

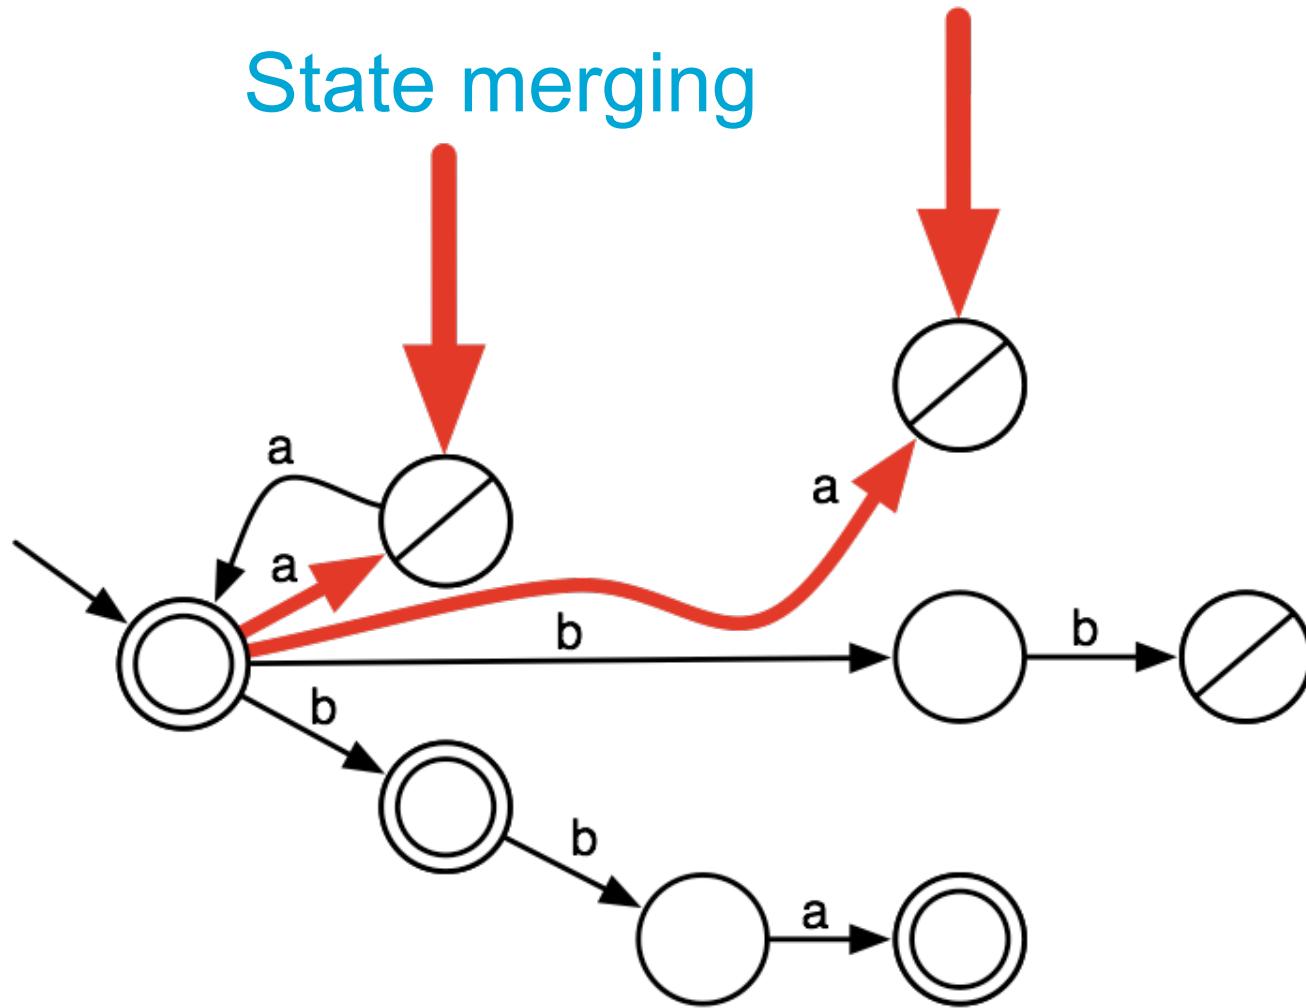
State merging



Determinization:

merge the targets of non-deterministic transitions

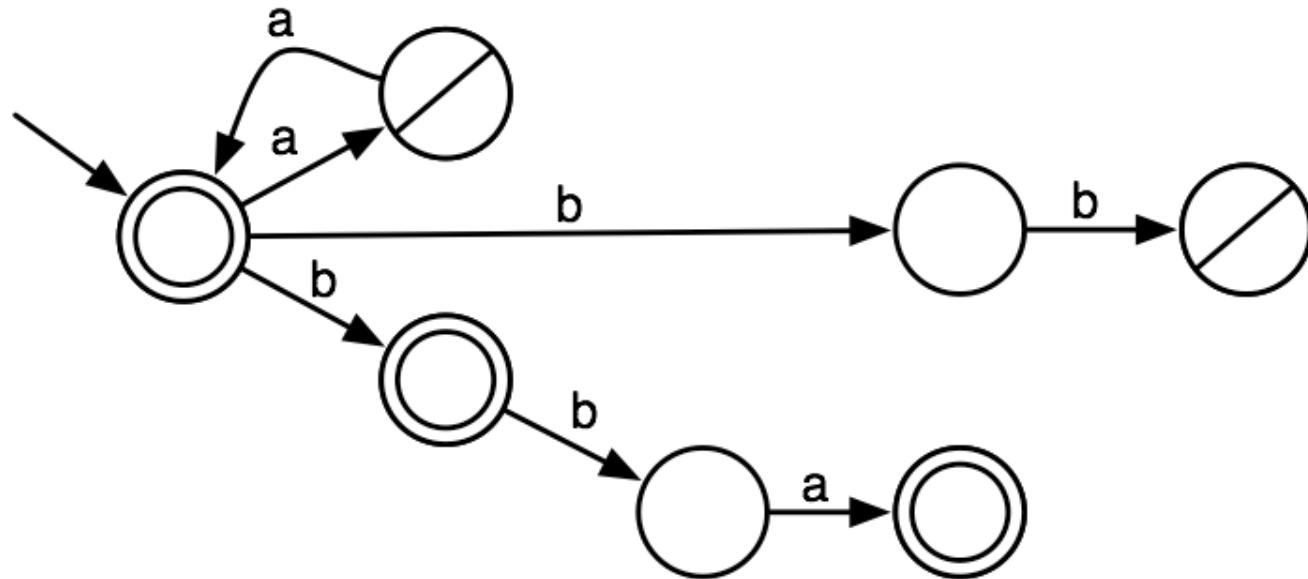
State merging



Determinization:

merge the targets of non-deterministic transitions

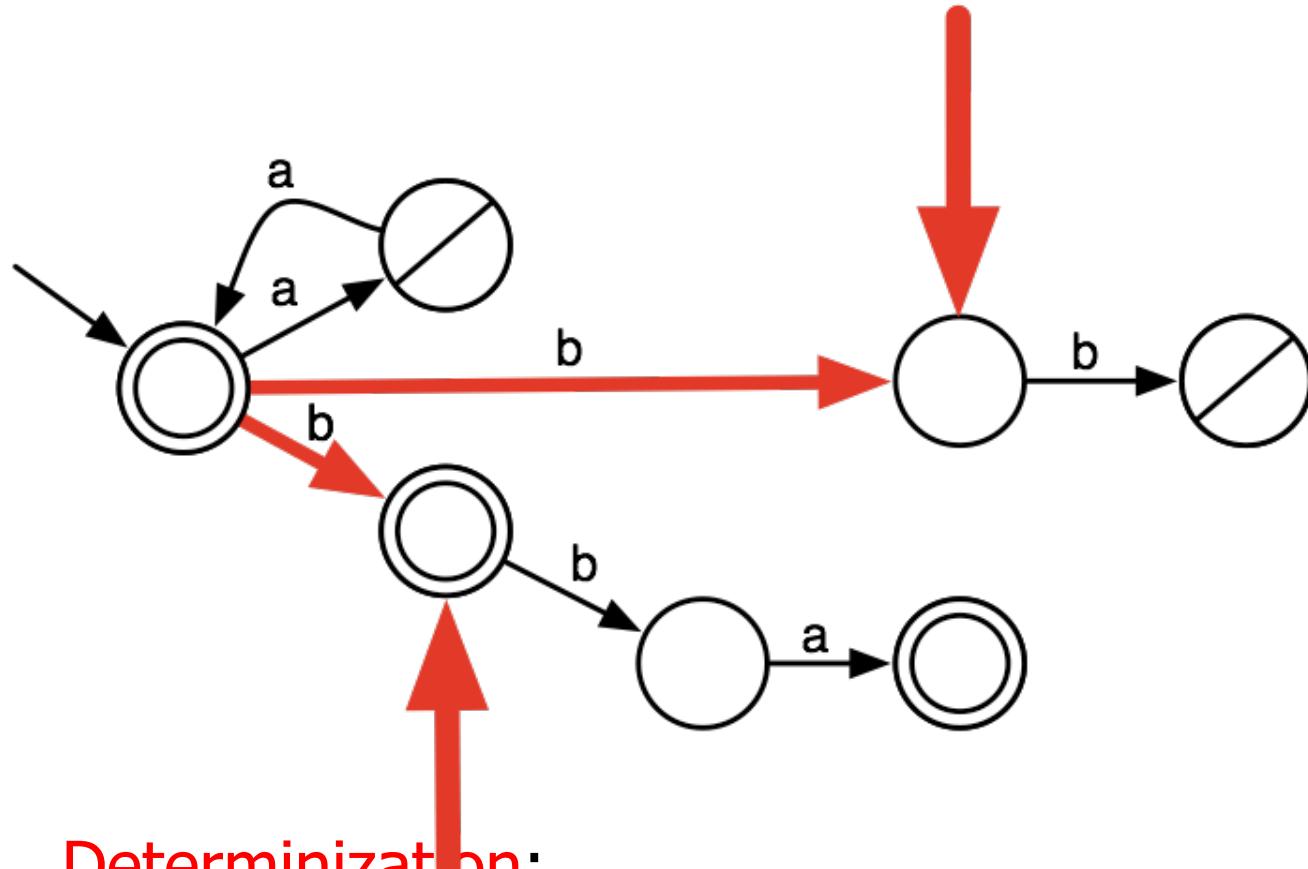
State merging



Determinization:

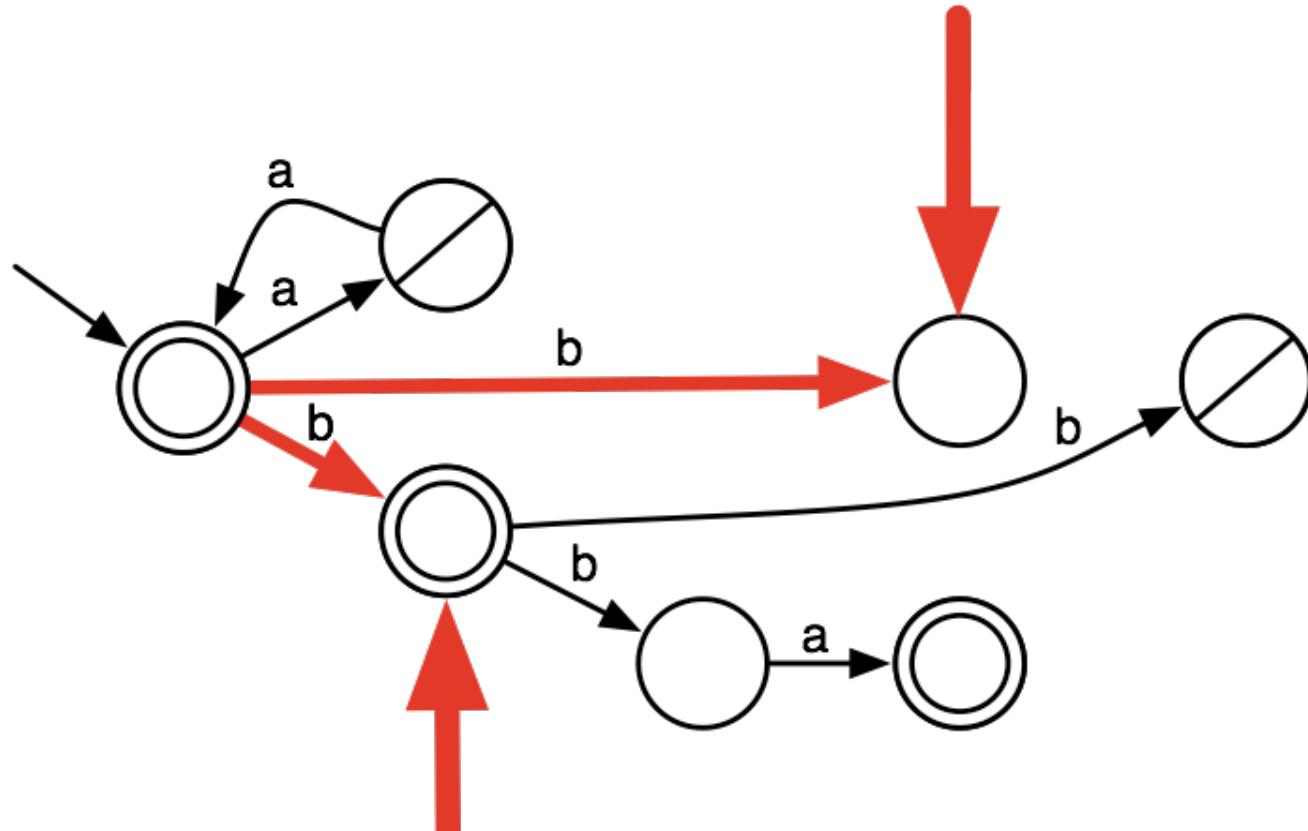
merge the targets of non-deterministic transitions

State merging



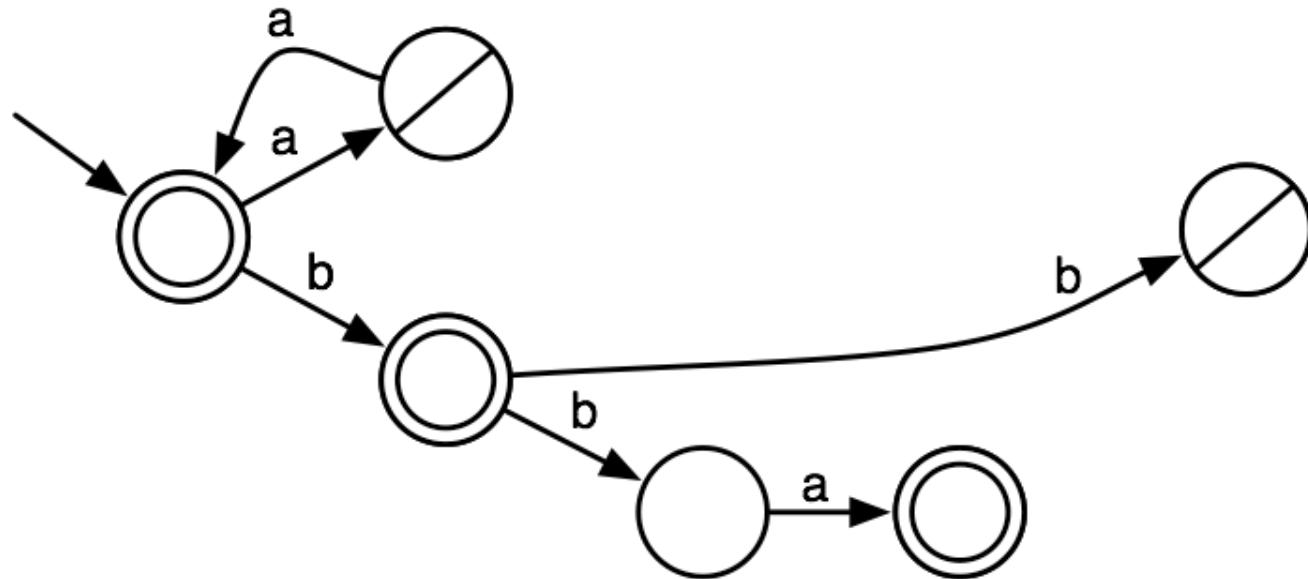
Determinization:
merge the targets of non-deterministic transitions

State merging



Determinization:
merge the targets of non-deterministic transitions

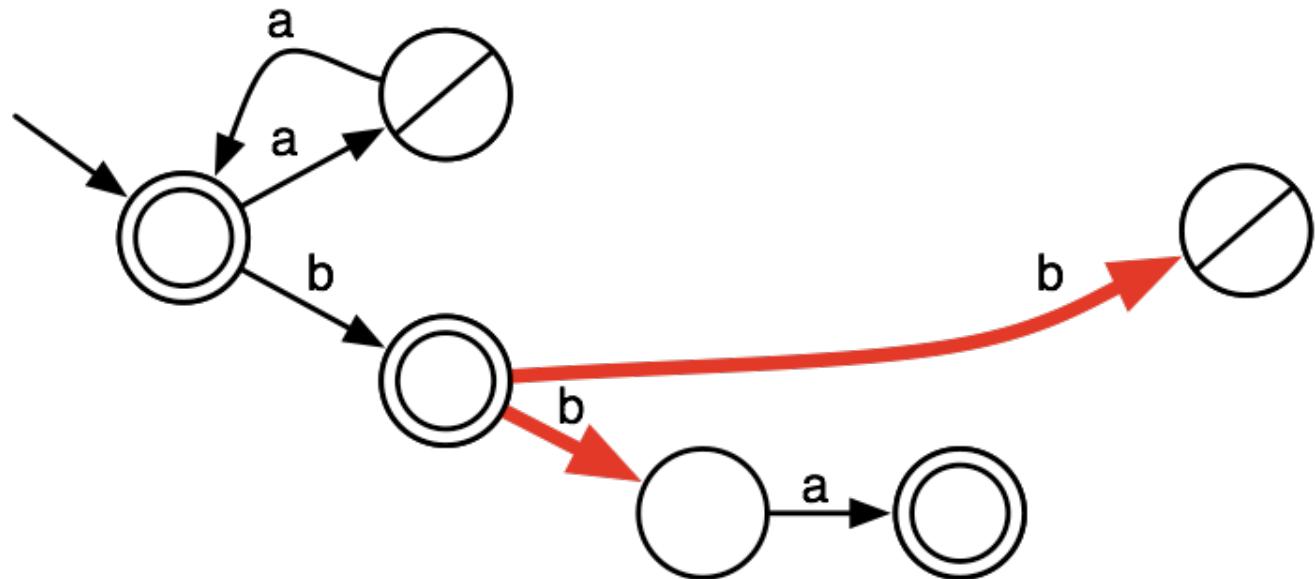
State merging



Determinization:

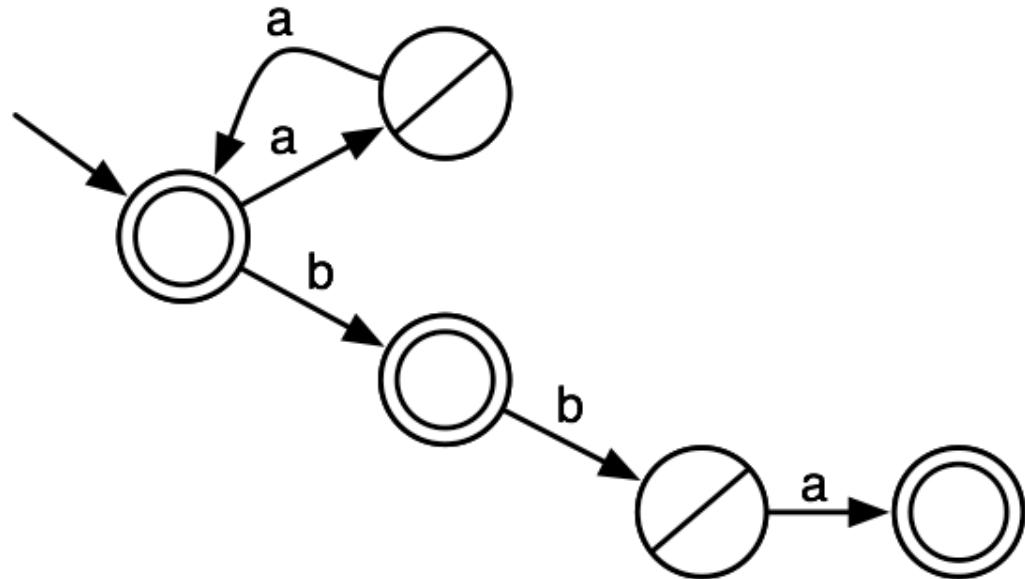
merge the targets of non-deterministic transitions

State merging



Select two new nodes to merge and **iterate**

State merging



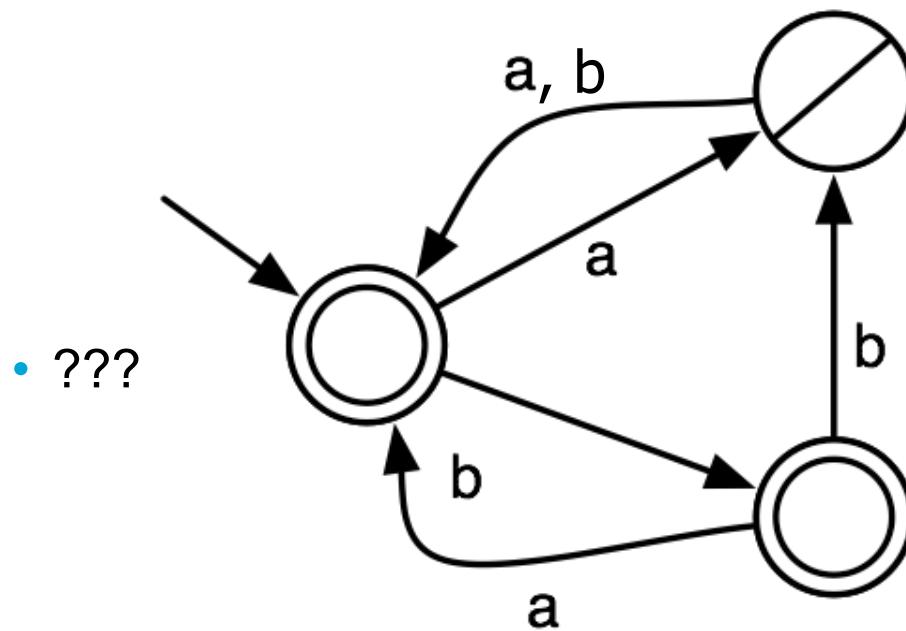
Select two new nodes to merge and **iterate**

State merging (RPNI)

1. Start from a **prefix tree acceptor**
2. Color the initial state red, and its children blue
3. Let B be the first blue state in **lexicographical order**
(a,b,aa,ab,ba,bb,aaa,aab,aba,abb,...)
4. For all red states R in lexicographical order
 1. **Merge B with R**
 2. Break if consistent, **undo** the merge otherwise
5. Color B red if it is not merged
6. Color all children of red states blue
7. Goto 3

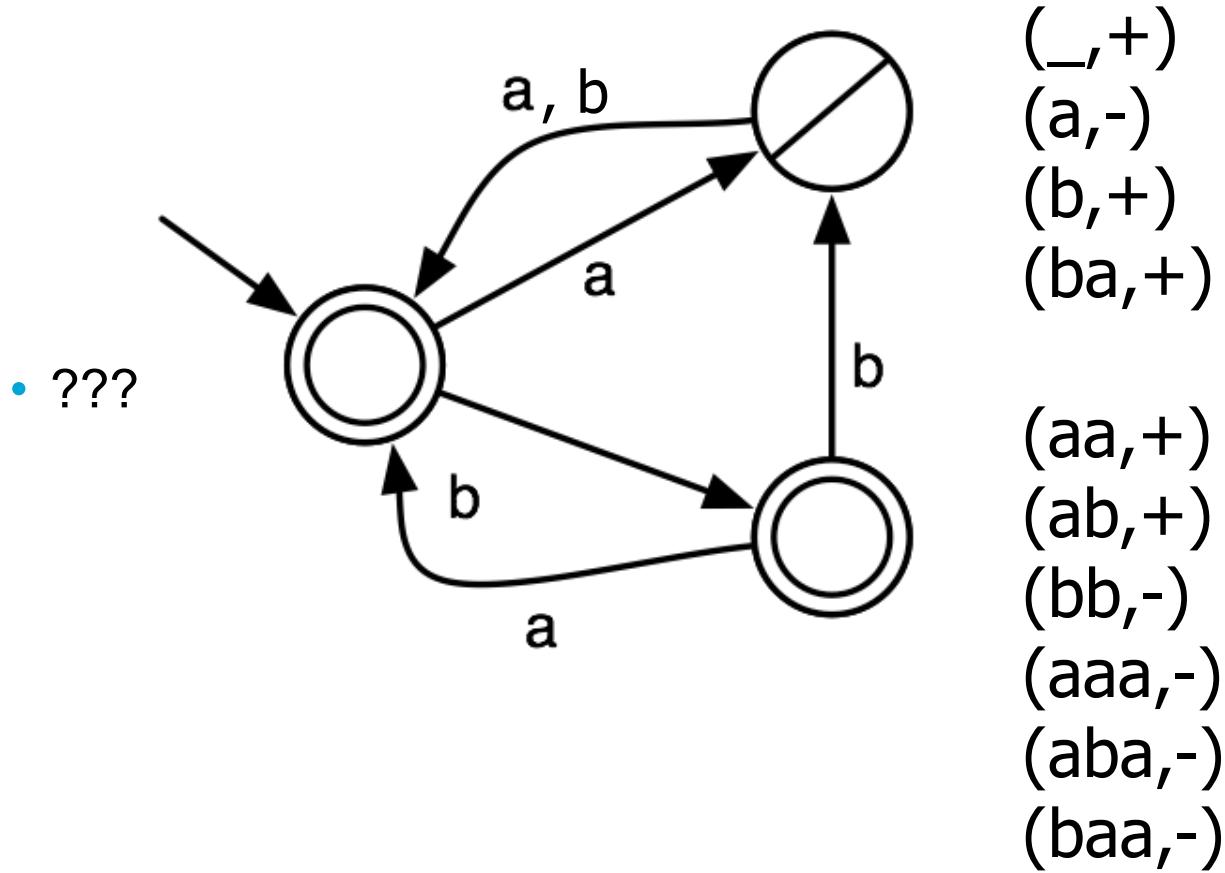
Convergence

- How much data does RPNI need to learn



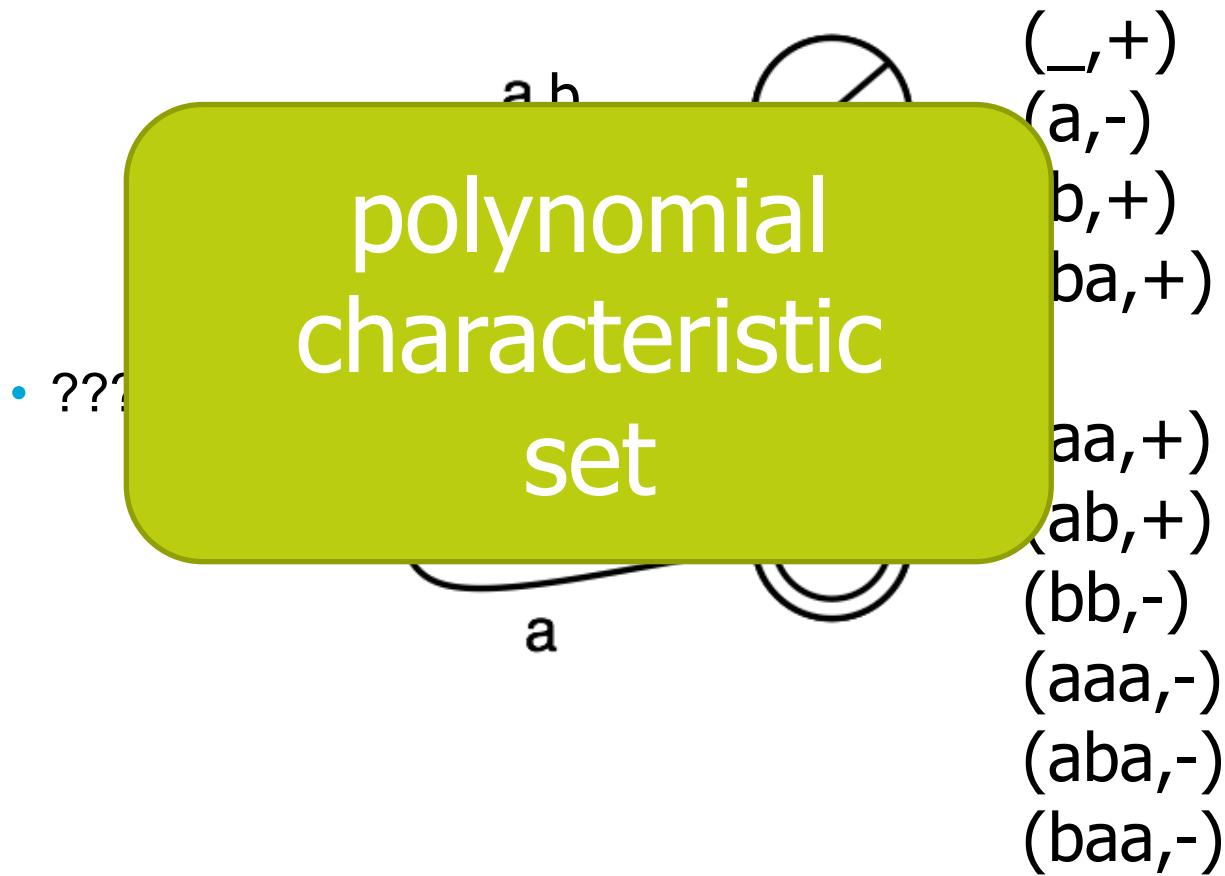
Convergence

- How much data does RPNI need to learn



Convergence

- How much data does RPNI need to learn



Results

- DFAs cannot be learned from unlabeled data
- Learning DFAs from labeled data is hard, but
 - In the limit, RPNI requires only a polynomial amount of data to converge
 - If the data source is nice, we only need a **polynomial amount of data and a polynomial-time algorithm**

State machine learning key ingredients

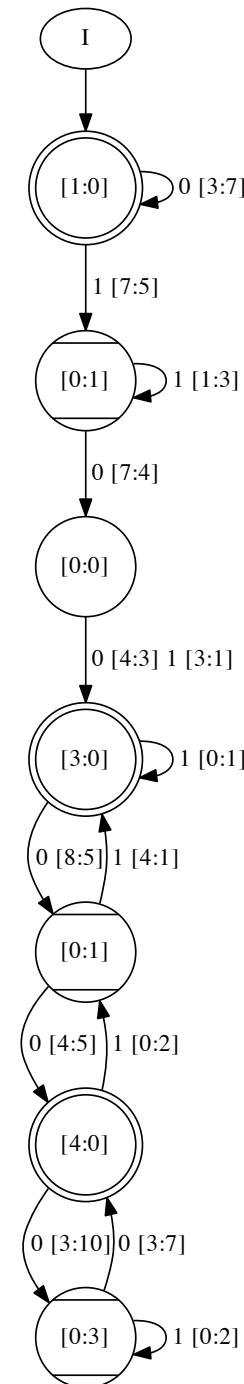
- When are two states **consistent**?
 - Positive gets merged with negative
 - Other events are possible
 - Statistics do not match
 - If the Markov property is violated
 - ...
- What to use as merge **score**?
 - Number of positive-positive, negative-negative merges
 - Overlap in sets of possible events
 - Statistical distance
 - ...

Passive automaton learning

- Learn a behavioral model from any event sequence
- Use the model as a classifier
- Analyze the model to understand its inner workings
- DFASAT, my own tool to learn these models
 - based on state merging
 - translates part of the learning problem to Satisfiability, and runs SAT-solvers to compute the best DFA
 - winner of the 2010 international state machine learning competition
 - allows many different consistency checks and score functions
 - possible to define your own
 - still needs work...
- Demo...

DFASAT produces

```
1910000000  
115011000101010100  
0501000  
1200  
012110000000100  
1510100  
13100  
06000001  
13101  
0201000001010000011100  
0  
0130111010000000  
13101  
171000000
```



DFASAT produces

19100000000
115011000101010100
0501000
1200
012110000000100
1510100
13100
06000001
13101
02010000010100000111000
0130111010000000
13101
171000000
...
10244 strings in total

